

Operating Systems: Lecture 6

Operating Systems: Lecture 6

Zombies

after `exit()`

Zombie

Calling `wait()` is important

Summary

Signals

Signals are software interrupts

How are signals generated?

Sending signals using the `kill` command

Foreground and background jobs

The `fg` command

The `kill()` system call

`SIGTSTP` and then `SIGTERM`

Some standard signals (man 7 signal)

Changing signal handlers

Example: change the behavior of Ctrl-C

Waiting for a signal

Timer

Summary

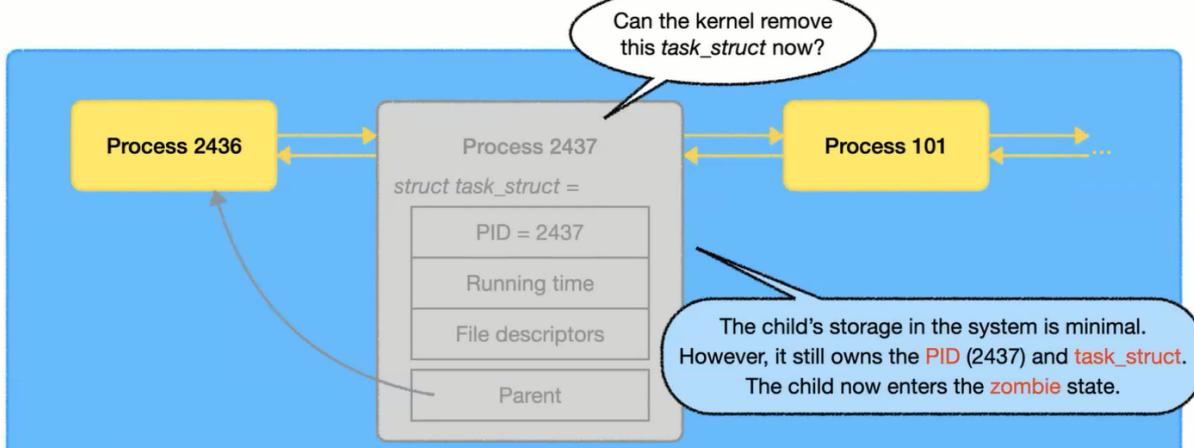
Midterm Exam

Zombies

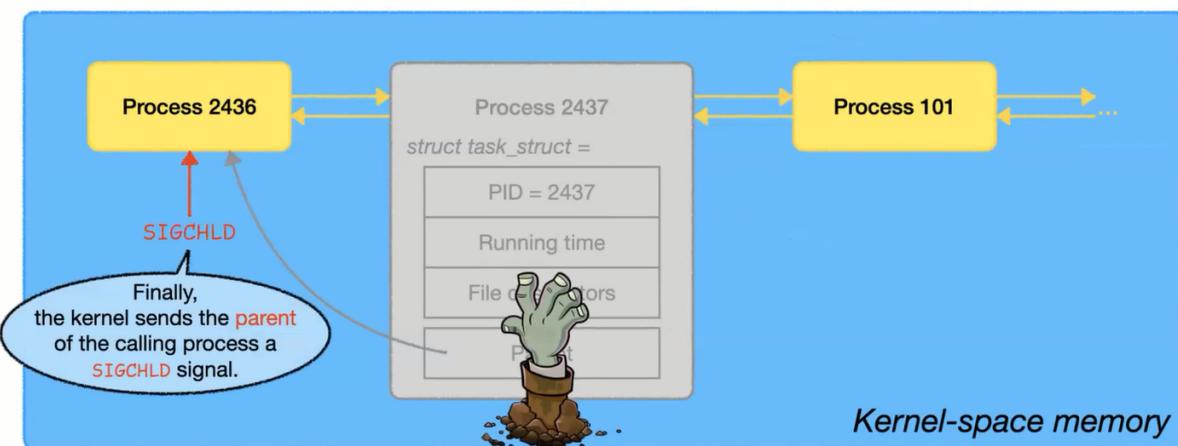
- after `exit()`

After `exit()` called, the kernel removes everything about the process, including the codes and files opened and other data. However, we shouldn't remove the task_struct, as the PID, status and exitcode are still needed by the parent process.

exit()



A `SIGCHLD` signal is sent to the parent process.

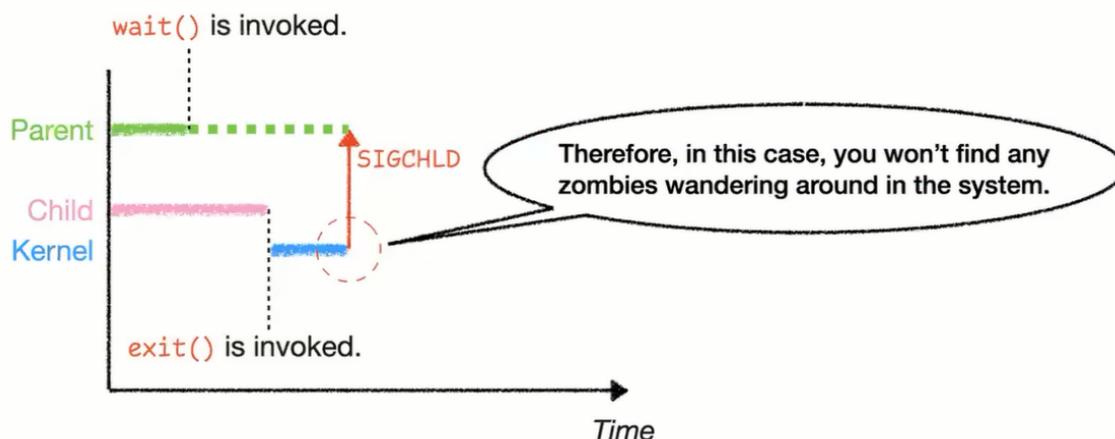


The `SIGCHLD` signal will be processed by the parent process, but it depends on the timing of `wait()`.

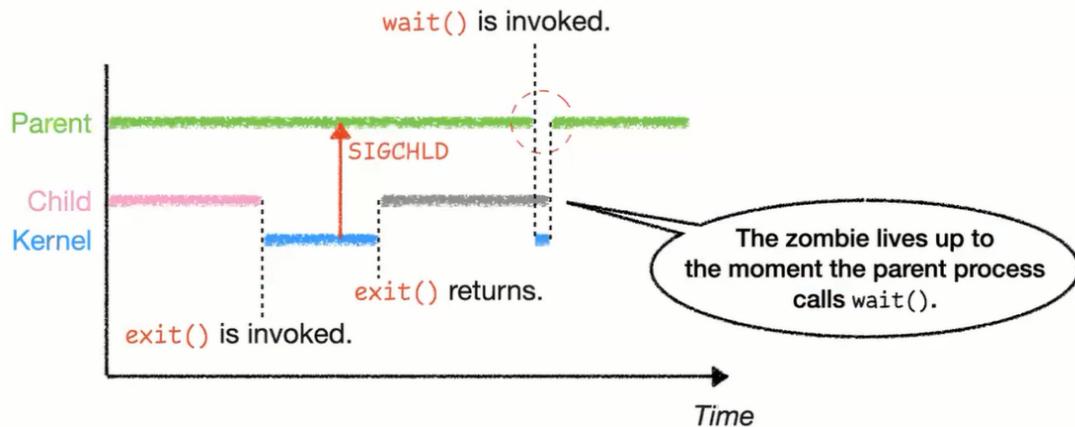
The `wait()` sets up a signal handler, which gets the exit code from the child and removes the `task_struct`.

There are two scenarios:

1. `wait()` called before `SIGCHLD`



2. `wait()` called after `SIGCHLD`



- Zombie

`wait()` and `waitpid()` are to reap zombie children.

You should never leave any zombies in the system.

Linux will label zombie processes as “`<defunct>`”.

```
int main() {
    pid_t pid;
    if ((pid = fork()) != 0) {
        printf("Look at Process %d...", pid);
        getchar(); // Press ENTER to continue
        wait(NULL);
        printf("Look again!");
        getchar(); // Press ENTER to continue
    }
}
```

zombie.c

Look at Process 20808...

Look again!

PID	TTY	STAT	TIME	COMMAND
19973	?	SN	0:00	sshd: yt2475@pts/9
19974	pts/9	SNs	0:00	-bash
20807	pts/9	SN+	0:00	./zombie
20808	pts/9	ZN+	0:00	[zombie] <defunct>
20839	?	SN	0:00	sshd: yt2475@pts/7
20840	pts/7	SNs	0:00	-bash
20973	pts/7	RN+	0:00	ps x
49141	?	SNsl	86:01	emacs --daemon

[yt2475@linserv1 ~]\$ ps x

PID	TTY	STAT	TIME	COMMAND
19973	?	SN	0:00	sshd: yt2475@pts/9
19974	pts/9	SNs	0:00	-bash
20807	pts/9	SN+	0:00	./zombie
20839	?	SN	0:00	sshd: yt2475@pts/7
20840	pts/7	SNs	0:00	-bash
21012	pts/7	RN+	0:00	ps x
49141	?	SNsl	86:01	emacs --daemon

"

`wait()` eliminated the zombie (defunct) from the system. The PID is removed from the processes.

- Calling `wait()` is important

Why should we remove all the zombies from the system?

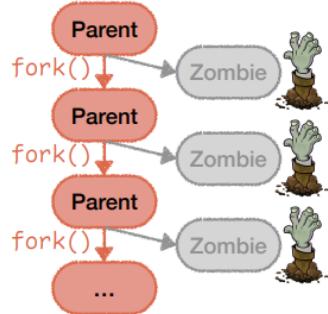
- We have already removed all the user-space memory and other kernel data. What it remains is just a PID and a node in the data struct. It takes only a tiny amount of memory.
- What could be the issue?
 - It takes a PID
- There's a maximum number of PID ($64k \approx 2^{16}$)

It's all about system resource management.

- A zombie takes up a PID.
- What would happen if we never clean up the zombies?

(Don't) try this...

```
int main() {  
    while (fork());  
}      bomb.c
```



```
$ cat /proc/sys/kernel/pid_max  
65536  
$ ./bomb  
-
```

```
$ ls  
No process left.  
$ poweroff  
No process left.  
$ _
```

- The code above calls `fork()` infinitely, so it will continuously create zombies and leave them in the system. The only active process running is the parent process.
- It will take up the maximum amount of PID
 - In this case it's 65536
- Then, you can't run any other process. You can't even turn off the computer, since you need `poweroff` to turn it off in Linux.

- Summary

It's important to know the internal workings of system calls.

Some system calls may produce surprising results.

Understanding those system calls helps you understand the process control in the operating system.

At least, you don't have to guess what would happen after invoking a system call.

Signals

- Signals are software interrupts

It's a form of **inter-process communication (IPC)**.

A process can send a **signal** to another process.

When a signal arrives, the OS interrupts the target process's normal control flow and executes the **signal handler**.

- In our previous example, the receiver is the parent process.
- After it receives the signal, its normal execution is interrupted, and the signal handler is invoked.

Example:

- What happens when you press *Ctrl-C* ?
 - It sends a **SIGINT** signal to the current process.
 - The default signal handler of **SIGINT** is to **terminate** the process.
 - The process can override this behavior by providing a signal handler.

```
[yt2475@linserv1 proc]$ gcc -o loop loop.c
[yt2475@linserv1 proc]$ ./loop
0
1
2
3
4
5
6
7
8
9
10
11
12
^C
```

- What about *Ctrl-Z* ?
 - It sends a **SIGTSTP** signal to stop the current process
 - Stop is different from terminating. It can be resumed by the signal **SIGCONT**

```

8
^Z
[1]+  Stopped                  ./loop
[yt2475@linserv1 proc]$ ps
  PID TTY          TIME CMD
19974 pts/9    00:00:00 bash
22156 pts/9    00:00:00 loop
22178 pts/9    00:00:00 ps
[yt2475@linserv1 proc]$ jobs
[1]+  Stopped                  ./loop
[yt2475@linserv1 proc]$ fg 1
./loop

```

- How are signals generated?

From the user space...

- Using **keyboard**: Ctrl-C sends **SIGINT**, Ctrl-Z sends **SIGTSTP** (terminal stop), Ctrl-\ sends **SIGQUIT**.
- Using **commands**: **kill**, **top** ...
- Using the **kill()** system call.

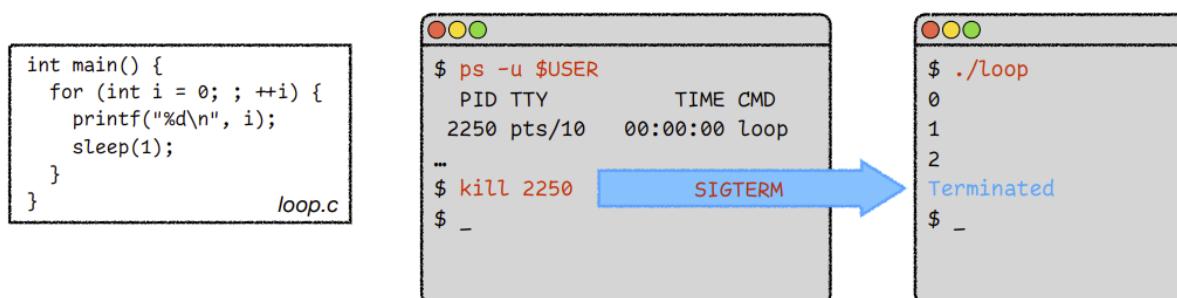
From the kernel or hardware...

- **SIGCHLD** (child stopped or terminated) comes from the kernel.
- **SIGFPE** (floating point exception, e.g., **division by zero**) comes from the CPU.
- **SIGSEGV** (invalid memory reference, a.k.a. **segmentation fault**) comes from the CPU to the kernel and then to the process.

- Sending signals using the **kill** command

The **kill** command sends **SIGTERM** to the target process by default.

The default signal handler of **SIGTERM** is to terminate the process



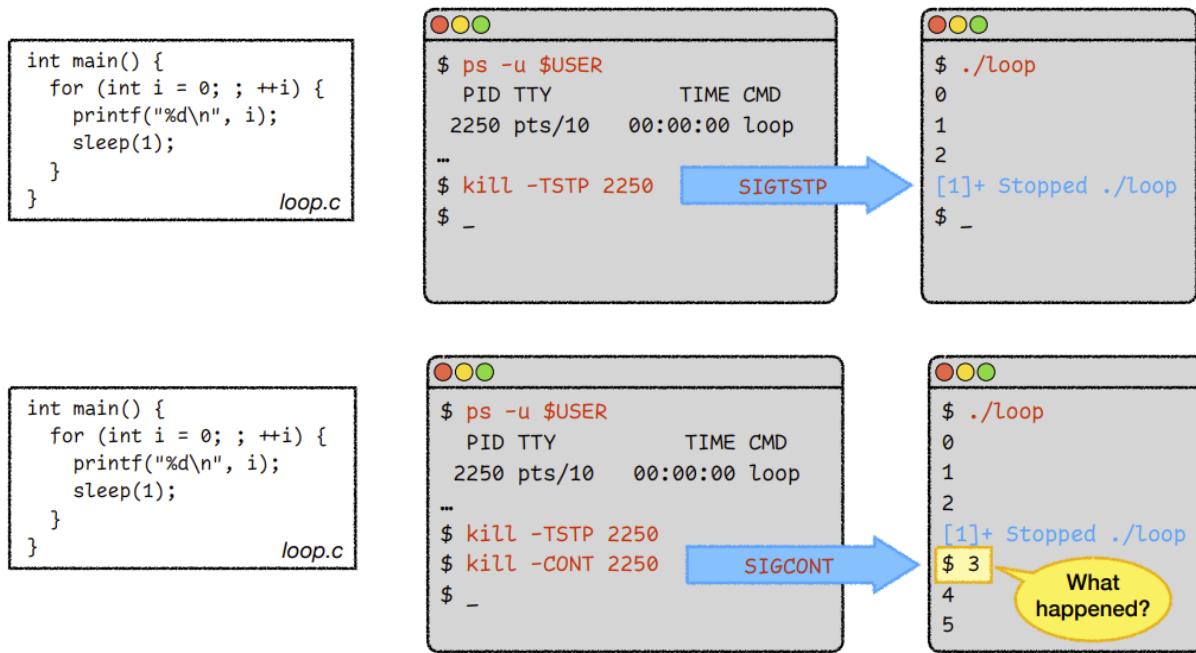
```

33
34
35
36
37
38
39
40
Terminated
[yt2475@linserv1 proc]$
  PID TTY      STAT   TIME COMMAND
19973 ?        SN      0:00 sshd: yt2475@pts/9
19974 pts/9    SNs    0:00 -bash
22300 pts/9    SN+    0:00 ./Loop
22322 ?        SN      0:00 sshd: yt2475@pts/7
22323 pts/7    SNs    0:00 -bash
22460 pts/7    RN+    0:00 ps x
49141 ?        SNSl   86:01 emacs --daemon
[yt2475@linserv1 ~]$ kill 22300
[yt2475@linserv1 ~]$

```

Actually, the `kill` command can send [any](#) signal.

Let's try `SIGTSTP` and `SIGCONT`.



- use `-TSTP` to send `SIGTSTP`
 - the same as `Ctrl-Z`

```

1
2
3
4
5
6
7
[1]+  Stopped                  ./loop
yt2475@linserv1 proc]$ 
    PID TTY      STAT   TIME COMMAND
 9973 ?        SN      0:00 sshd: yt2475@pts/9
 9974 pts/9    SNs     0:00 -bash
 2322 ?        SN      0:00 sshd: yt2475@pts/7
 2323 pts/7    SNs     0:00 -bash
 2986 pts/9    SN+    0:00 ./loop
 2990 pts/7    RN+    0:00 ps x
 9141 ?        SNSl   86:01 emacs --daemon
yt2475@linserv1 ~]$ kill -TSTP 22986

```

- Use `-CONT` to send `SIGCONT`

```

26
27
[1]+  Stopped                  ./loop
[yt2475@linserv1 proc]$ 28
29
30
31
32

19973 ?        SN      0:00 sshd: yt2475@pts/9
19974 pts/9    SNS     0:00 -bash
22322 ?        SN      0:00 sshd: yt2475@pts/7
22323 pts/7    SNS     0:00 -bash
22986 pts/9    SN+    0:00 ./loop
22990 pts/7    RN+    0:00 ps x
49141 ?        SNSl   86:01 emacs --daemon
[yt2475@linserv1 ~]$ kill -TSTP 22986
[yt2475@linserv1 ~]$ kill -CONT 22986

```

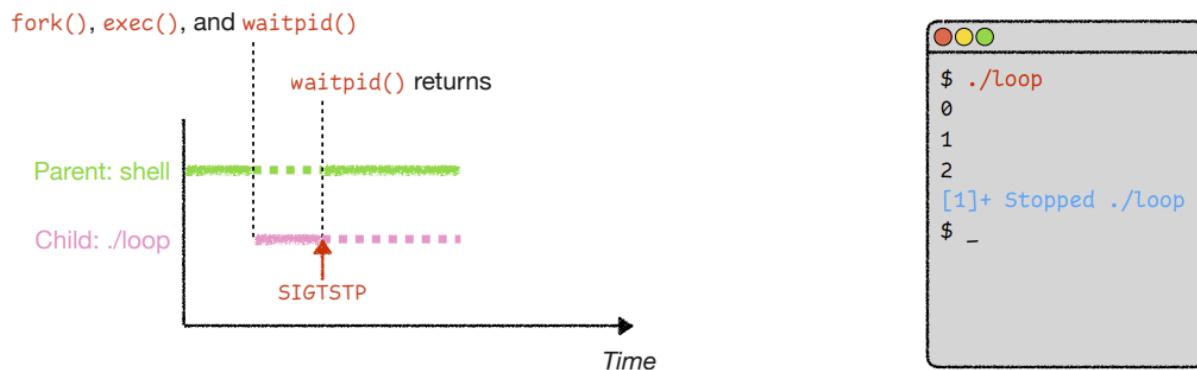
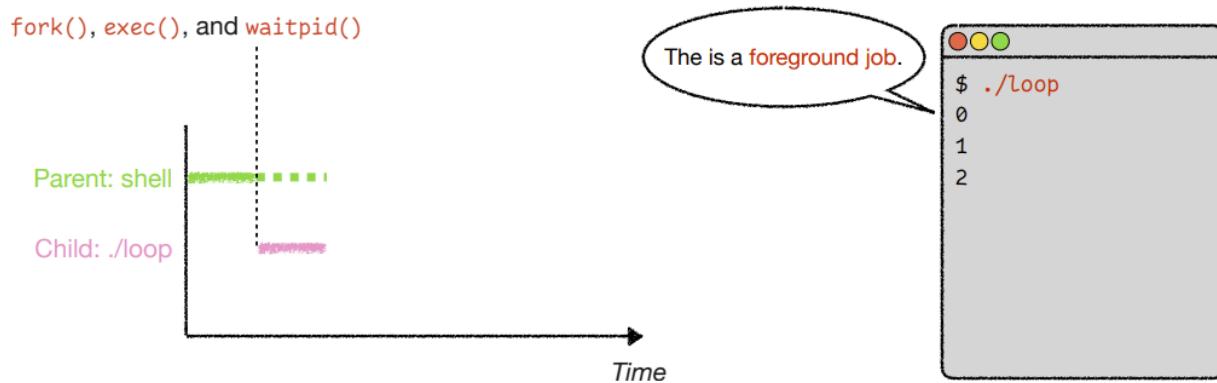
- The `kill` command is actually a command to send signal. Although it by default kills a process, the true purpose of the command is to send signals
- However, after the process is resumed by `SIGCONT`, it can't be stopped by `Ctrl-C` !
 - Why?

```

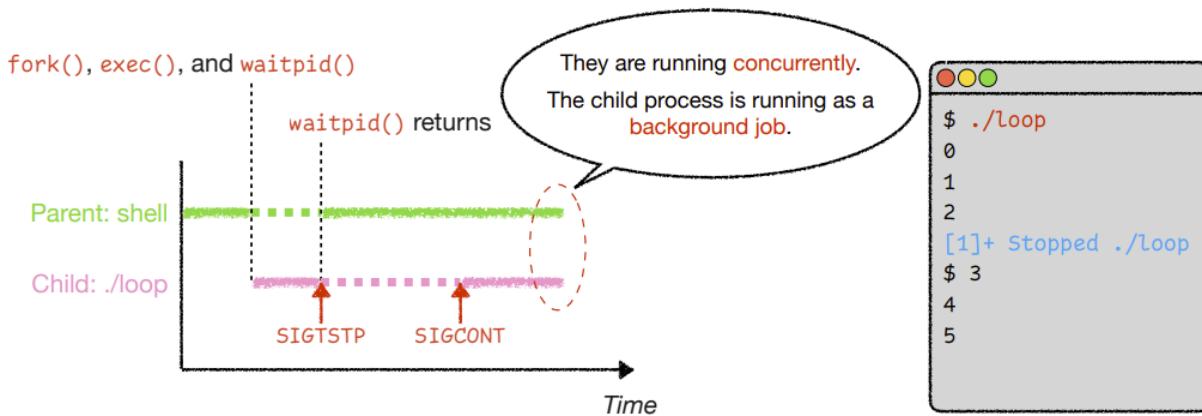
103
104
^C
[yt2475@linserv1 proc]$ 105
106
107
^C
[yt2475@linserv1 proc]$ 108
109

```

- Foreground and background jobs

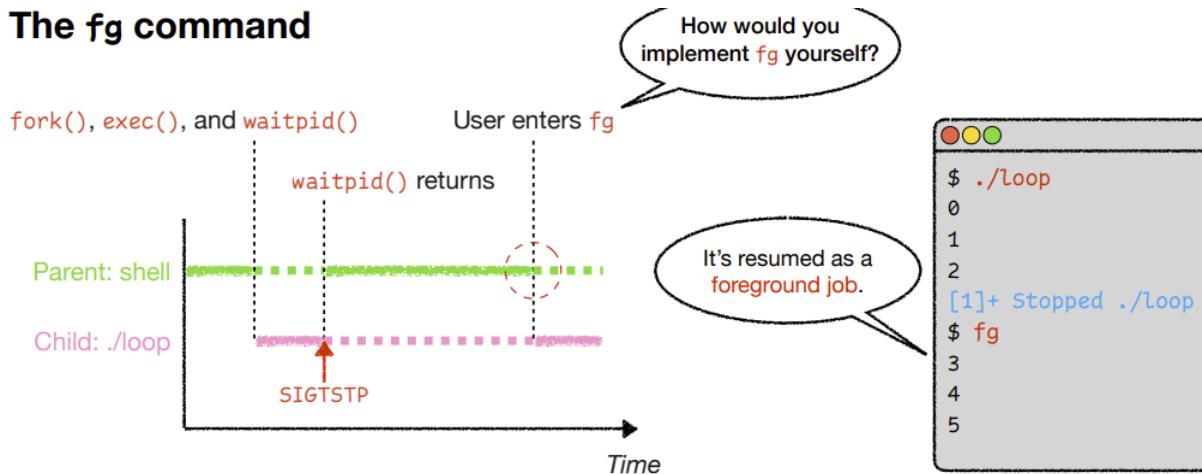


- Here, the shell calls `wait()`, and it will be blocked until the loop (child process) terminates or stops
- As long as the foreground program is still running, the shell will block itself. It won't take any command until the child stops or terminates



- Here, we send a `SIGCONT` to the process from another window, and it continues running
- However, the shell **has no idea** that the process has continued running. So the shell is not blocking itself anymore.
 - As a result, **Both** the shell and the process are running at the same time.
- Note the shell runs as the **foreground job**, and the child process runs as the **background job**.
 - When you press Ctrl-C or other keys in shell, the signal is only sent to the shell, not the child process
 - Only shell has access to your keyboard input
 - This is not a good thing
 - Therefore, we should use the `fg` command to resume a job.

• The `fg` command



- After you used the `fg` command, the shell knows that you want to resume that process.
 - **The shell** will send a `SIGCONT` to the process
 - it'll block itself to wait for this process
- The difference is that `fg` is initiated by the **Shell**
 - So the shell knows you want to continue this process. It'll send a `SIGCONT` and wait
 - If you send `SIGCONT` from another process instead of the shell, it has no idea that you have resumed the process. So it has no idea that it needs to wait.

- To stop the loop process before, you can type `fg` in shell.
 - The shell will regain control of the child process.
 - Now you can use `Ctrl-C` to stop the process

- The `kill()` system call

Just like the `kill` command, the `kill()` syscall **sends a signal** to a process.

```
int kill(pid_t pid, int sig);
```

The `raise()` library function sends a signal to yourself.

```
int raise(int sig);
```

How would you implement `raise()`?

- To implement `raise()`, you can just use `getPID()` and `kill()`

- **SIGTSTOP** and then **SIGTERM**

- What would happen if we stop a process and then terminate it?

```
20

[1]+  Stopped                  ./loop
[yt2475@linserv1 proc]$
[yt2475@linserv1 proc]$
[yt2475@linserv1 proc]$
[yt2475@linserv1 proc]$ fg
./loop
Terminated
[yt2475@linserv1 proc]$

19973 ?          SN      0:00 sshd: yt2475@pts/9
19974 pts/9    SNs     0:00 -bash
22322 ?          SN      0:00 sshd: yt2475@pts/7
22323 pts/7    SNs     0:00 -bash
23957 pts/9    SN+     0:00 ./loop
23966 pts/7    RN+     0:00 ps x
49141 ?          SNSl   86:01 emacs --daemon
[yt2475@linserv1 ~]$ kill -TSTP 23957
[yt2475@linserv1 ~]$ kill -TERM 23957
```

- What we did is that we sent a `SIGTSTOP` to the process to stop it.
- Then we send a `SIGTERM` to it when it's stopped, trying to terminate that process
 - Because the process was stopped, it couldn't handle the `SIGTERM`.

- So it seems like nothing happens.
- The `SIGTERM` is a pending signal. It will be handled whenever the process is resumed
- When we use `fg`, the process is continued
- Right after it's continued, the `SIGTERM` is handled, and the process is terminated
- However, we can also send a `SIGKILL` to kill the process immediately

```

11
12
13
14

[1]+  Stopped                  ./loop
[yt2475@linserv1 proc]$
[1]+  Killed                   ./loop
[yt2475@linserv1 proc]$

19973 ?          SN      0:00 sshd: yt2475@pts/9
19974 pts/9      SNs     0:00 -bash
22322 ?          SN      0:00 sshd: yt2475@pts/7
22323 pts/7      SNs     0:00 -bash
24182 pts/9      SN+    0:00 ./loop
24189 pts/7      RN+    0:00 ps x
49141 ?          SNSl   86:01 emacs —daemon
[yt2475@linserv1 ~]$ kill -TSTP 24182
[yt2475@linserv1 ~]$ kill -KILL 24182

```

- Some standard signals (man 7 signal)

Signal	Default handler	Default handler
SIGINT	Interrupt from keyboard (Ctrl-C).	Terminate the process.
SIGTERM	Termination (the default signal sent by the <i>kill</i> command).	Terminate the process.
SIGKILL	Kill, no matter what.	Terminate the process (cannot be overridden).
SIGTSTP	Stop typed at terminal (Ctrl-Z).	Stop the process.
SIGCONT	Continue if stopped.	Continue the process.
SIGCHLD	Child stopped or terminated.	Ignore the signal.
SIGFPE	Floating point exception.	Terminate the process and dump core.
SIGSEGV	Invalid memory reference (segmentation fault).	Terminate the process and dump core.

- **Changing signal handlers**

The `signal()` system call changes the **signal handler**.

```
void ( *signal(int sig, void (*handler)(int)) ) (int);
```

I know...I know it's hard to read. Let's do some `typedef`.

```
typedef void (*sighandler_t)(int);
sighandler_t signal(int sig, sighandler_t handler);
```

Returns the **old signal handler**.

The **signal number**.

A **function pointer** to the **new signal handler**, which takes an **int** argument and returns **void**.

- Example: change the behavior of Ctrl-C

```

void handler(int sig) {
    printf("...Ouch!\n");
}

int main() {
    signal(SIGINT, handler);
    for (int i = 0; ; ++i) {
        printf("%d\n", i);
        sleep(1);
    }
}                                     sig.c

```

The new signal handler.

Register a new signal handler when Ctrl-C is pressed.

When Ctrl-C is pressed, the process is interrupted, and the new signal handler is invoked.

How to terminate this program?

A terminal window with a red close button, a yellow minimize button, and a green maximize button. The command \$./sig is entered, followed by several numbers (0, 1, 2, 3, 4, 5, ...) each preceded by a blue '^C' character, indicating Ctrl-C interruptions. The final output is '...Ouch!'.

```

$ ./sig
0
1
2
^C...Ouch!
3
4
^C...Ouch!
5
...

```

```

1 #include <signal.h>
2 #include <stdio.h>
3 #include <unistd.h>
4
5 void handler (int sig) {
6     printf(" ... Ouch!\n");
7 }
8
9 int main() {
10     signal(SIGINT, handler);
11     int i;
12     for (i=0; ;++i) {
13         printf("%d\n", i);
14         sleep(1);
15     }
16 }

```

```

26
27
^C...Ouch!
28
29
^C^C...Ouch!
30
^C...Ouch!
31

```

```

53
...Ouch!
54
55
56
...Ouch!
57
58

19973 ? SN 0:00 sshd: yt2475@pts/9
19974 pts/9 SNs 0:00 -bash
22322 ? SN 0:00 sshd: yt2475@pts/7
22323 pts/7 SNs 0:00 -bash
25043 pts/9 SN+ 0:00 ./sig
25081 pts/7 RN+ 0:00 ps x
49141 ? SNSl 86:01 emacs —daemon
[yt2475@linserv1 ~]$ kill -INT 25043
[yt2475@linserv1 ~]$ kill -INT 25043
[yt2475@linserv1 ~]$

```

- We changed the default signal handler for the signal `SIGINT`, so when we press Ctrl-C or use `kill` to send the signal, it'll only print Ouch, but the process won't be stopped
- Now, we can still use `SIGTSTOP` to stop the process and use `SIGTERM` to terminate the process. However, we can also overwrite these two handlers.
- The only signal we can't overwrite is `SIGKILL`. If we overwrite every other signal handlers, we could only use `SIGKILL` to kill the process

“

There are also some signal reserved for user to define, like `SIGUSR1`, `SIGUSR2`.

But there's a fixed amount of signals

• Waiting for a signal

The `pause()` system call puts the process to **block (sleep)** until the delivery of...

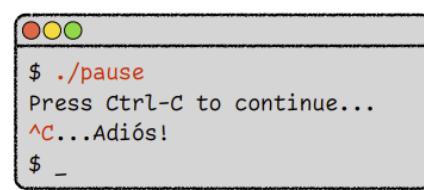
- a signal **handled** by the process, or
- a signal **terminating** the process.

```

void handler(int sig) {
}

int main() {
    signal(SIGINT, handler);
    printf("Press Ctrl-C to continue...\n");
    pause();
    printf("...Adiós!\n");
}

```



- **Timer**

The `alarm()` system call sets up a **one-time asynchronous timer** for the process.

- A ***SIGALRM*** signal will be sent to the process at the timeout.
- The default ***SIGALRM*** handler terminates the process, but you can override it.

The `setitimer()` system call sets up an **interval timer**.

Can you implement `sleep()`?

- **Summary**

Signal is a kind of **software interrupt**.

It has many quirky behaviors.

`kill()` is not intended to kill anybody, but to **send signals**.

When in doubt, refer to the **man pages**.

Midterm Exam

Announcement

Mid-term exam

Due to scheduling difficulties, our **mid-term exam** will be held **online** via **Brightspace Quizzes**.

Time & venue: Thursday, March 9, 12:30 p.m.–1:45 p.m. @ anywhere.

You must **connect to the Zoom session** and **turn on your webcam** for the entire duration of the exam.

As a result, our mid-term exam will be open-book.

Our final exam will still be in-person and closed-book (with one “cheat sheet”).

You can refer to...

- Slides.
- Textbooks.
- Docker & CIMS servers (e.g., read man pages, use gcc to try out programs).
- Online websites (you can search existing information, but **you are not allowed to ask questions and seek answers**).

However, keep in mind that by doing these things, you will likely **waste a lot of time** with little gain.

You must not communicate with any other people or AI-based assistants.