

OS: Lecture 9



OS: Lecture 9

Review

Scheduling Algorithms

Trade-offs

Priority scheduling

Static priority scheduling

Limitations

Dynamic priority scheduling

Can we do better?

Multilevel feedback queue (MLFQ)

Summary

Interprocess communication

What is IPC?

Why do we need IPC?

How to do IPC?

Case study: pipe

Shared memory

What could go wrong?

Out of sync?

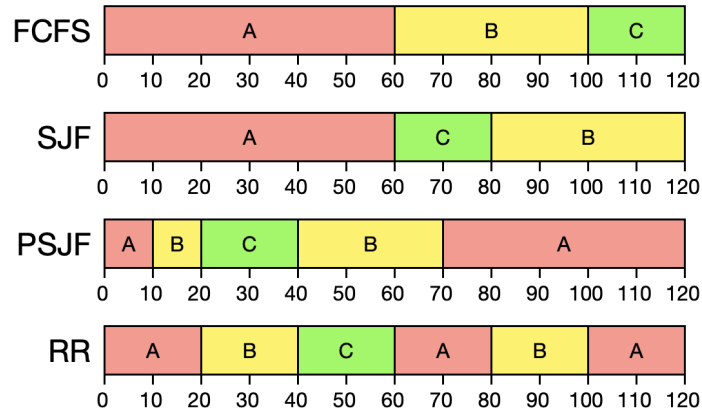
Race conditions

Review

Scheduling algorithms

FCFS, SJF, PSJF, RR

Job	Arrival time	CPU requirement
A	0	60
B	10	40
C	20	20



Scheduling Algorithms

- Trade-offs

This is an inherent **trade-off** between **performance** and **fairness**.

A **fair** scheduler (such as RR) evenly divides the CPU among active jobs on a small time scale, at the cost of turnaround time.


Most ordinary users run a lot of **interactive jobs** on modern operating systems.

- They value **responsiveness** more than CPU **efficiency**.



Another trade-off comes from **context switching**. It's relatively slow.

Case 1: time slice = 10ms, context switch = 1ms. (~10% of time is wasted.)



Case 2: time slice = 100ms, context switch = 1ms. (<1% of time is wasted.)



- Short time slice \Rightarrow many context switches \Rightarrow low CPU efficiency.
- Long time slice \Rightarrow poor response \Rightarrow “sluggish.”

- Priority scheduling

Each job is assigned a **priority**.

The scheduler always chooses the job with the **highest priority** to run.


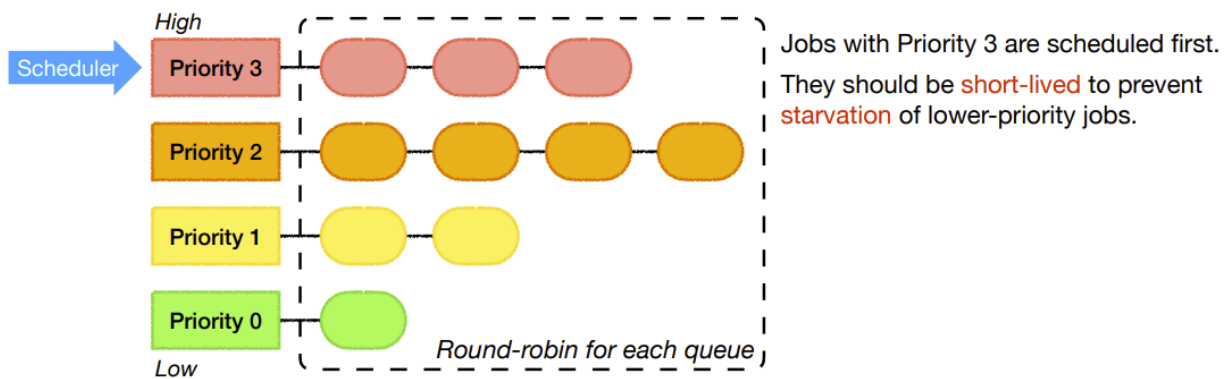
Priorities can be **static** or **dynamic**.


Static priority means that a job is assigned a **fixed** priority when it is submitted to the system.

- **Example:** a background email process should get a lower priority than a real-time video game process.

Dynamic priority means that a job's priority may be **changing** throughout its life in the system.

- Static priority scheduling





- Limitations

Limitations

High-priority jobs may run for a prolonged period, or even indefinitely.

Low-priority jobs may **starve** to death.

- **Rumor:** when the IBM 7094 mainframe at MIT was shut down in 1973, people found a low-priority process submitted in 1967 had not yet been run.

It does not differentiate between CPU-bound and I/O-bound jobs.

- I/O-bound jobs spend most of their time waiting for I/O to complete.
- When such a job wants the CPU, we'd better schedule it immediately to let it start its next I/O request.
- In that way, I/O requests can proceed **in parallel** with another process actually computing.

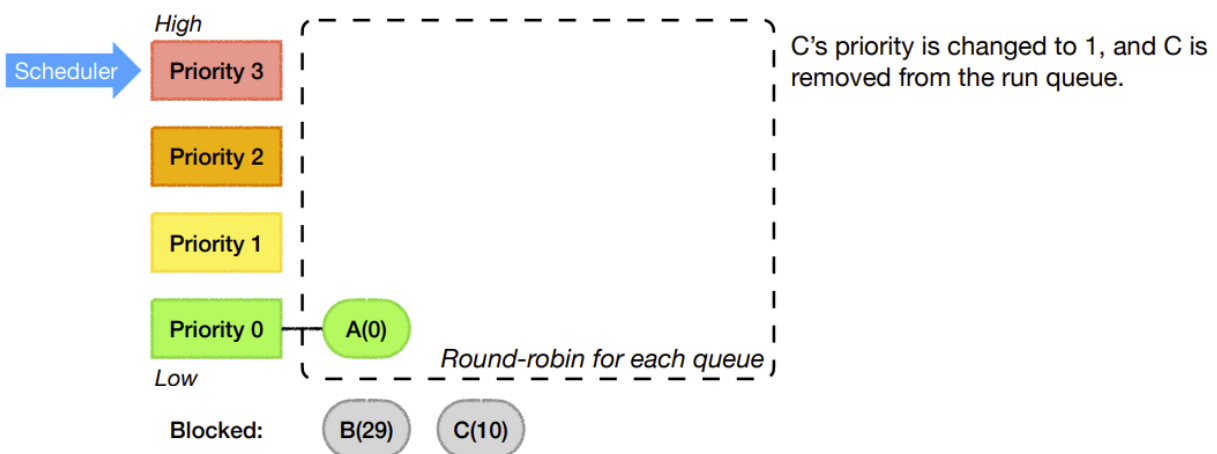
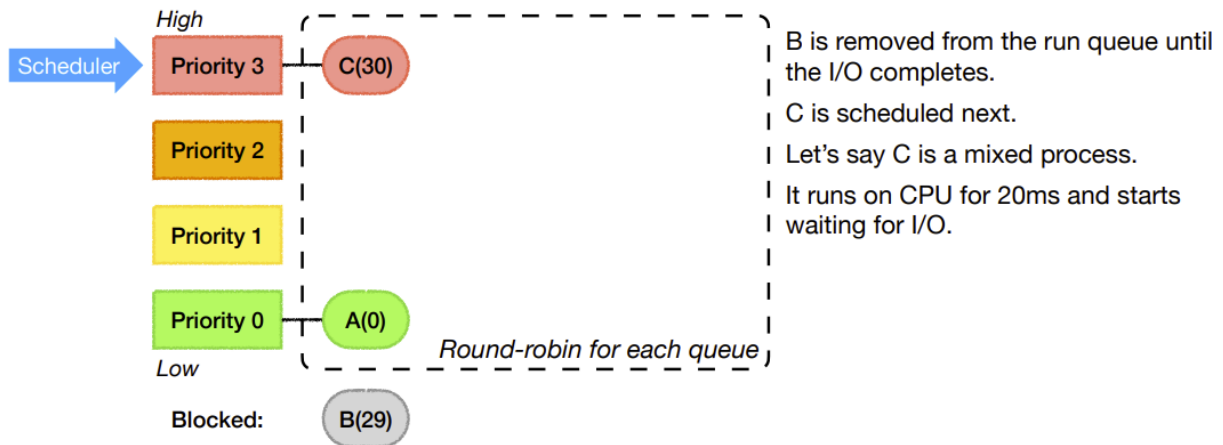
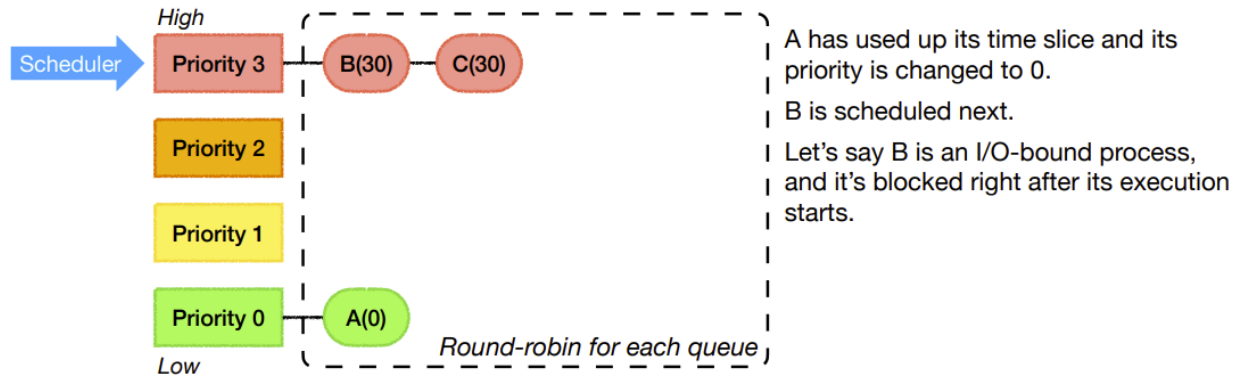
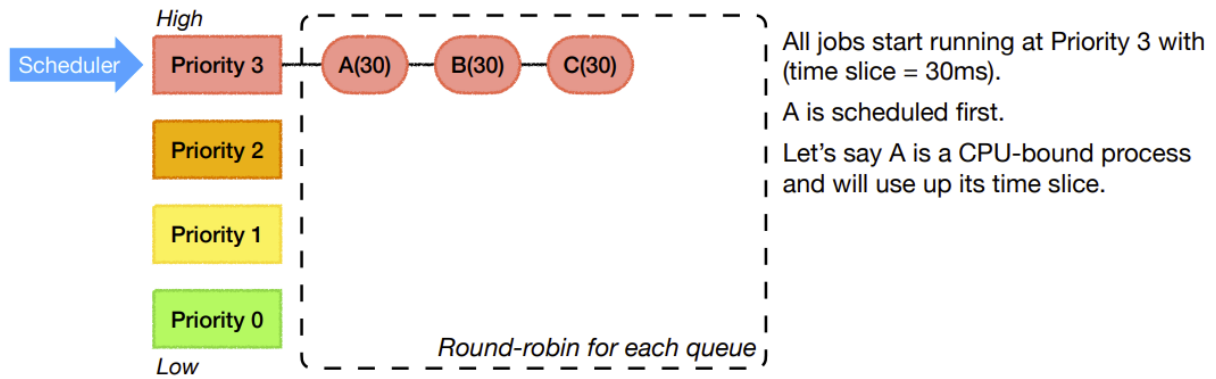
- Dynamic priority scheduling

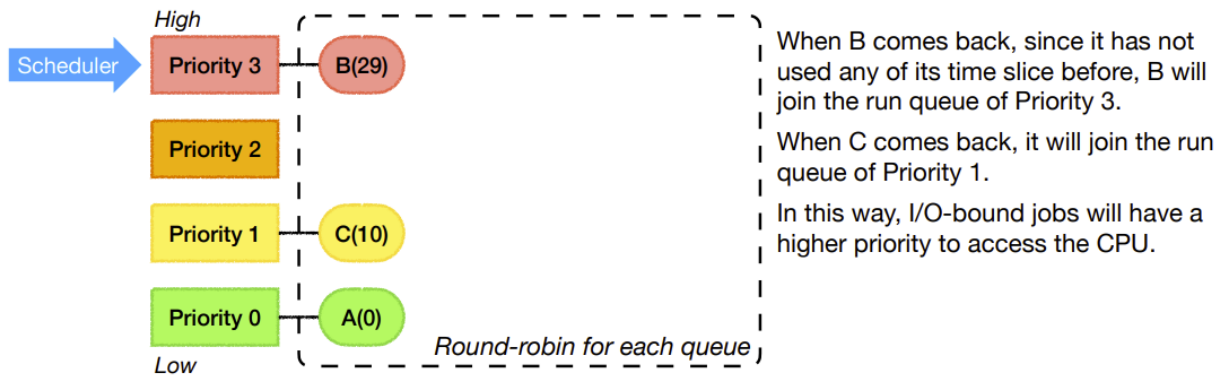
There is **no standard way** to assign priorities dynamically.

Let's look at an **example** policy.

Rules:

- All jobs start running at Priority 3 with time slice = 30ms.
- A job is preempted if its time slice is used up or it starts waiting for I/O.
- When a job is preempted, its priority is changed to $\left\lceil \frac{\text{its time slice left}}{10\text{ms}} \right\rceil$.





– Can we do better?

Remember the **trade-off** between **efficiency** and **responsiveness**?

Ideally, we want to...

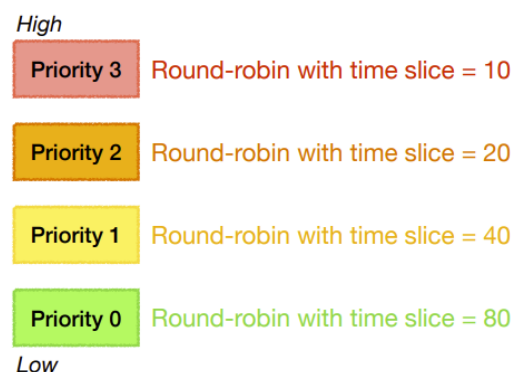
- Optimize **turnaround** time.
 - Run shorter jobs first.
 - Give CPU-bound jobs a large time slice to reduce context switching.
- Make the system feel **responsive** to interactive users.
 - Can't give all jobs a large time slice.
 - Minimize response time.

• Multilevel feedback queue (MLFQ)

It's a kind of dynamic priority scheduling, but **each priority has its own policy**.

Rules:

- All jobs start running at the highest priority.
- When a job uses up its time slice, its priority is reduced by 1.
- If a job gives up the CPU before the time slice is up, it stays at the same priority.




“

This is just an example


In reality, we might have way more levels than it. We could have 20 levels or even over a hundred.

Example 1: a single long-running job




- Long-running jobs get a large time slice to reduce context switching. ✓

Example 2: along came a short job



- Short jobs run first. ✓

Example 3: what about I/O?



- I/O-bound jobs have a higher priority to access the CPU. ✓

Are there any limitations?

- Long-running jobs may **starve** if there are *too many interactive jobs*.
- What if a program **changes its behavior** over time?
 - A job may need to run for a long time when it first starts. After that, it becomes interactive.
 - Such a job will be punished forever.

New rule: the priority boost

- After some time period, move all jobs to the highest priority.

Here's a summary of what we've achieved so far...

Rule 1: If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).

Rule 2: If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in round-robin fashion using the time slice of the given queue.

Rule 3: When a job enters the system, it is placed at the highest priority (the topmost queue).

Rule 4: Once a job uses up its time slice at a given level, its priority is reduced (it moves down one queue).

Rule 5: After some time period, move all the jobs in the system to the topmost queue.

MLFQ observes how jobs behave over time, and prioritize them accordingly.

- It can deliver excellent overall **performance** (similar to SJF/PSJF) for short-running interactive jobs.
- It's **fair** and makes progress for long-running CPU-intensive workloads.

Therefore, many modern operating systems use a form of MLFQ as their base scheduler.

- Summary

So, *what's the best scheduling algorithm?*

Unfortunately, there is **no best or standard** algorithm, partly because...

- We can't predict the CPU requirement of a process.
 - We don't even know if a job will eventually terminate!
- Online scheduling is an NP-hard problem.

Linux employs the **Completely Fair Scheduler (CFS)** since kernel 2.6.23.

- It's a **dynamic priority scheduling** algorithm based on **red-black trees**.

Interprocess communication

- What is IPC?

Processes often need to **communicate** with one another.

This is called **interprocess communication (IPC)**.

Have you used any methods of IPC?

- Signal: `kill -STOP 2250`
- Pipeline: `ls | less`
- File, socket, shared memory...

- Why do we need IPC?

To share information

- Of course...

To reuse software


- You can implement a spell checker by...

```
curl "https://en.wikipedia.org/wiki/Pipeline_(Unix)" | sed 's/^[^a-zA-Z ]/ /g' | tr 'A-Z' 'a-z\n' |  
grep '[a-z]' | sort -u | comm -23 - <(sort /usr/share/dict/words) | less
```


To speedup computation

- **Example:** MapReduce.
- You can divide a job into tasks, run them as various processes **in parallel**, and combine the results.
- To learn more about MapReduce and big data analytics systems, take CSCI-UA.0476 or CSCI-GA.2436.

- How to do IPC?



Message passing:



- Case study: pipe

Example: `ls | less`

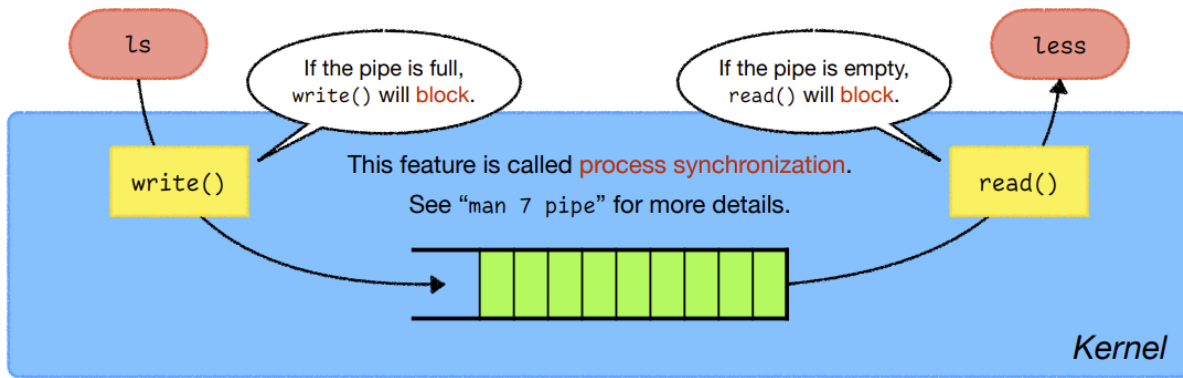
From a programmer's perspective...

The `pipe()` system call returns two file descriptors: `pipefd[0]` and `pipefd[1]`.



This is called a **producer-consumer model**.

From the kernel's perspective...



- Shared memory

Shared memory is...

- A region of memory created by the kernel;
 - See "[man 7 shm_overview](#)" for more details.
- Visible to all processes in the system;
 - By contrast, a pipe is only visible to two processes.
- Accessible by all processes in the system.
 - However, there are syscalls to change the ownership and permissions of the shared memory.

- What could go wrong?

- In the case of a **pipe**, the **kernel** provides a form of **synchronization**.
- However, for **shared memory**, it's up to the **processes** to coordinate.
- What could possibly go wrong?



- Out of sync?

Out of sync?

```
int main() {  
    // create a 4-byte shared memory  
    int *count = mmap(NULL, 4, PROT_READ | PROT_WRITE,  
                      MAP_SHARED | MAP_ANONYMOUS, -1, 0);  
  
    pid_t pid = fork();  
  
    for (int i = 0; i < 1000000; ++i) {  
        ++*count;  
    }  
  
    if (pid) {  
        wait(NULL);  
        printf("count = %d\n", *count);  
    }  
}
```

What does it output?

shm.c



```
$ ./shm  
count = 1434181  
$ ./shm  
count = 1062142  
$ ./shm  
count = 1256405  
$ _
```

“

`mmap()` system call can create a shared memory

```
int main() {  
    // create a 4-byte shared memory  
    int *count = mmap(NULL, 4, PROT_READ | PROT_WRITE, MAP_SHARED | MAP_ANONYMO  
  
    pid_t pid = fork();  
  
    for (int i = 0; i < 1000000; ++i) {  
        ++*count;  
    }  
  
    if (pid) {  
        wait(NULL);  
        printf("count = %d\n", *count);  
    }  
}
```

```
[yt2475@linserve1 proc]$ gcc -o shm shm.c  
[yt2475@linserve1 proc]$ ./shm  
count = 1157264  
[yt2475@linserve1 proc]$ ./shm  
count = 1251921  
[yt2475@linserve1 proc]$ ./shm  
count = 1190847
```

```

int main() {
    // create a 4-byte shared memory
    int *count = mmap(NULL, 4, PROT_READ | PROT_WRITE,
                      MAP_SHARED | MAP_ANONYMOUS, -1, 0);

    pid_t pid = fork();

    for (int i = 0; i < 1000000; ++i) {
        ++*count;
    }

    if (pid) {
        wait(NULL);
        printf("count = %d\n", *count);
    }
}
shm.c

```

```

$ gcc -S shm.c
$ cat shm.s
...
    movq    -16(%rbp), %rax
    movl    (%rax), %eax
    leal    1(%rax), %edx
    movq    -16(%rbp), %rax
    movl    %edx, (%rax)
...

```

```

$ gcc -S shm.c
$ cat shm.s
...
    movq    -16(%rbp), %rax
    movl    (%rax), %eax
    leal    1(%rax), %edx
    movq    -16(%rbp), %rax
    movl    %edx, (%rax)
...

```

a = count; → a = *count;
 a = *a;
 d = a + 1; → d = a + 1;
 a = count;
 *a = d; → *count = d;

What could go wrong?

	Parent	Child	count
Parent is running	a = *count; (a = 0)		0
	d = a + 1; (d = 1)		0
Context switch	-----		
Child is running		a = *count; (a = 0)	0
		d = a + 1; (d = 1)	0
		*count = d; (*count = 1)	1
Context switch	-----		
Parent is running	*count = d; (*count = 1)		1



– Race conditions

This scenario is called a **race condition** (or, more specifically, a **data race**).

The results depend on the **timing** of the execution, i.e., the particular **order** in which the **shared resource** is accessed.

Race conditions are always bad...

- Worse yet, **compiler optimizations** may generate crazy output if your code has data races.
- What if you compile the previous code with “gcc -O1” and “gcc -O2”?
- To learn more about **undefined behavior** (a.k.a. “**nasal demons**”), read **Schrödinger’s Code**.

Because the computation is **nondeterministic**, debugging is no fun at all.

- Heisenbug: bugs that disappear or change behavior when you try to debug.

“

This is also one of the **undefined behaviors** in C

The **undefined behaviors** come from data races.

If your code contains a data race, it's undefined, and the compile can choose whatever code it want to generate.

Will cause Heisenbug:

- it seems to work when you try to debug

```
[yt2475@linserv1 proc]$ gcc -O1 -o shm shm.c
[yt2475@linserv1 proc]$ ./shm
count = 1000000
[yt2475@linserv1 proc]$ ./shm
count = 1000000
[yt2475@linserv1 proc]$ ./shm
count = 1000000
[yt2475@linserv1 proc]$ ./shm
count = 1000000
```

```
[yt2475@linserv1 proc]$ gcc -O2 -o shm shm.c
[yt2475@linserv1 proc]$ ./shm
count = 2000000
[yt2475@linserv1 proc]$ ./shm
count = 2000000
[yt2475@linserv1 proc]$ ./shm
count = 2000000
[yt2475@linserv1 proc]$ ./shm
count = 2000000
```