

# OS Lecture 10

## OS Lecture 10

Recall: race conditions (Data Races)

Mutual exclusion

Critical sections

Requirements

Attempt #1: disabling interrupts

Attempt #2: using a “lock” variable

Attempt #3: strict alternation

Attempt #4: Peterson’s algorithm

Attempt #5: spinlocks

Attempt #5.1: too much spinning?

Attempt #6: semaphores

Summary

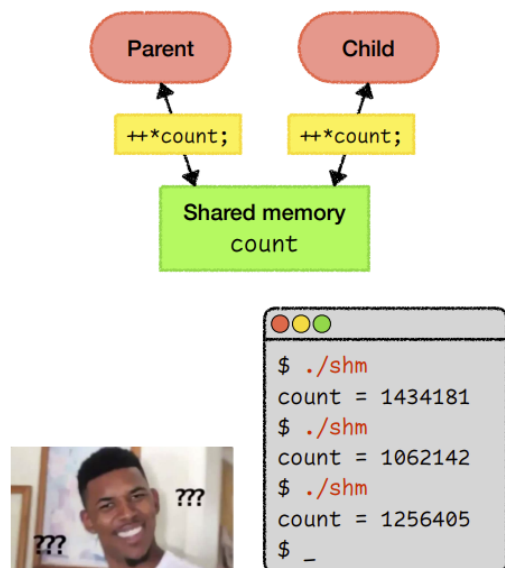
## Recall: race conditions (Data Races)

### Out of sync?

```
int main() {  
    // create a 4-byte shared memory  
    int *count = mmap(NULL, 4, PROT_READ | PROT_WRITE,  
                      MAP_SHARED | MAP_ANONYMOUS, -1, 0);  
  
    pid_t pid = fork();  
  
    for (int i = 0; i < 1000000; ++i) {  
        ++*count;  
    }  
  
    if (pid) {  
        wait(NULL);  
        printf("count = %d\n", *count);  
    }  
}
```

What does it output?

shm.c



```

int main() {
    // create a 4-byte shared memory
    int *count = mmap(NULL, 4, PROT_READ | PROT_WRITE,
                      MAP_SHARED | MAP_ANONYMOUS, -1, 0);

    pid_t pid = fork();

    for (int i = 0; i < 1000000; ++i) {
        ++*count;
    }

    if (pid) {
        wait(NULL);
        printf("count = %d\n", *count);
    }
}

```

shm.c

```

$ gcc -S shm.c
$ cat shm.s
...
    movq    -16(%rbp), %rax
    movl    (%rax), %eax
    leal    1(%rax), %edx
    movq    -16(%rbp), %rax
    movl    %edx, (%rax)
...

```

```


$ gcc -S shm.c
$ cat shm.s
...
    movq    -16(%rbp), %rax
    movl    (%rax), %eax
    leal    1(%rax), %edx
    movq    -16(%rbp), %rax
    movl    %edx, (%rax)
...

```

What could go wrong?

a = count;	→	a = *count;
a = *a;	→	a = *count;
d = a + 1;	→	d = a + 1;
a = count;	→	*count = d;
*a = d;	→	*count = d;

	Parent	Child	count
Parent is running	a = *count; (a = 0)		0
	d = a + 1; (d = 1)		0
Context switch	-----		
Child is running		a = *count; (a = 0)	0
		d = a + 1; (d = 1)	0
		*count = d; (*count = 1)	1
Context switch	-----		
Parent is running	*count = d; (*count = 1)		1



This scenario is called a **race condition** (or, more specifically, a **data race**).

The results depend on the **timing** of the execution, i.e., the particular **order** in which the **shared resource** is accessed.

Race conditions are always bad...

- Worse yet, **compiler optimizations** may generate crazy output if your code has data races.
- What if you compile the previous code with “gcc -O1” and “gcc -O2”?
- To learn more about **undefined behavior** (a.k.a. “**nasal demons**”), read **Schrödinger’s Code**.

Because the computation is **nondeterministic**, debugging is no fun at all.

- **Heisenbug**: bugs that disappear or change behavior when you try to debug.

- **Mutual exclusion**

To avoid race conditions, we need **mutual exclusion**:

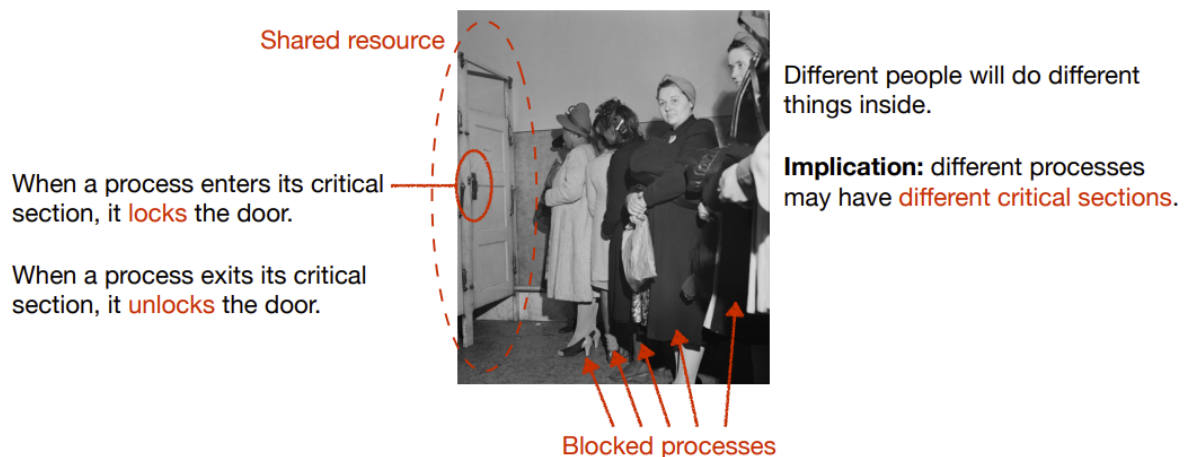
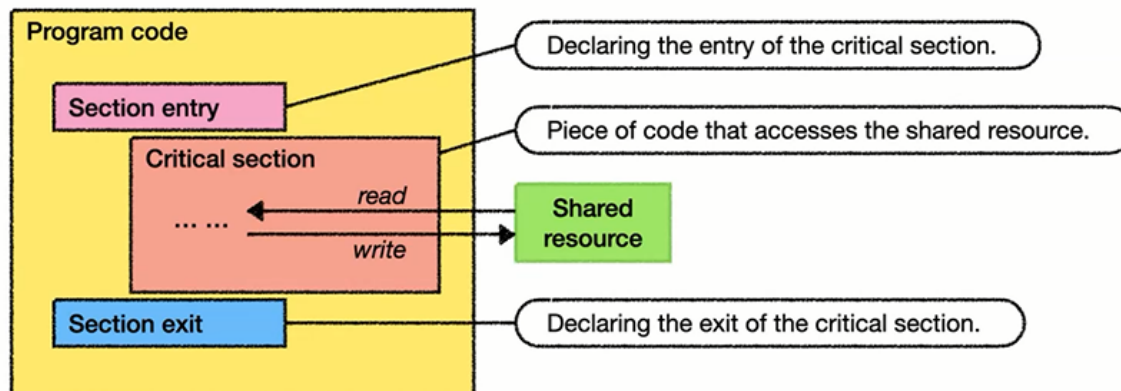
- If one process is accessing a shared resource,
- The other processes must be **excluded** from accessing the same thing.

**Note:** **race condition** is a **problem**, and **mutual exclusion** is a **requirement to avoid such a problem**.

- However, mutual exclusion may **hinder the performance of parallel computations**.

## Critical sections

A **critical section** (a.k.a. **critical region**) is a **piece of code** that accesses a shared resource.




A critical section is a **piece of code**, not a shared resource.

A critical section should be **as tight as possible**.

- What would happen if you declare the **entire program** as a big critical section?

A critical section may access **multiple shared resources**.

Mutual exclusion is required if the intersection is not empty.



## • Requirements

**No two processes may be simultaneously inside their critical sections.**

- This is the **mutual exclusion** requirement: when one process is inside its critical section, any attempts to go inside the critical sections by other processes are **not allowed**.

**No assumptions may be made about the speeds or the number of CPUs.**

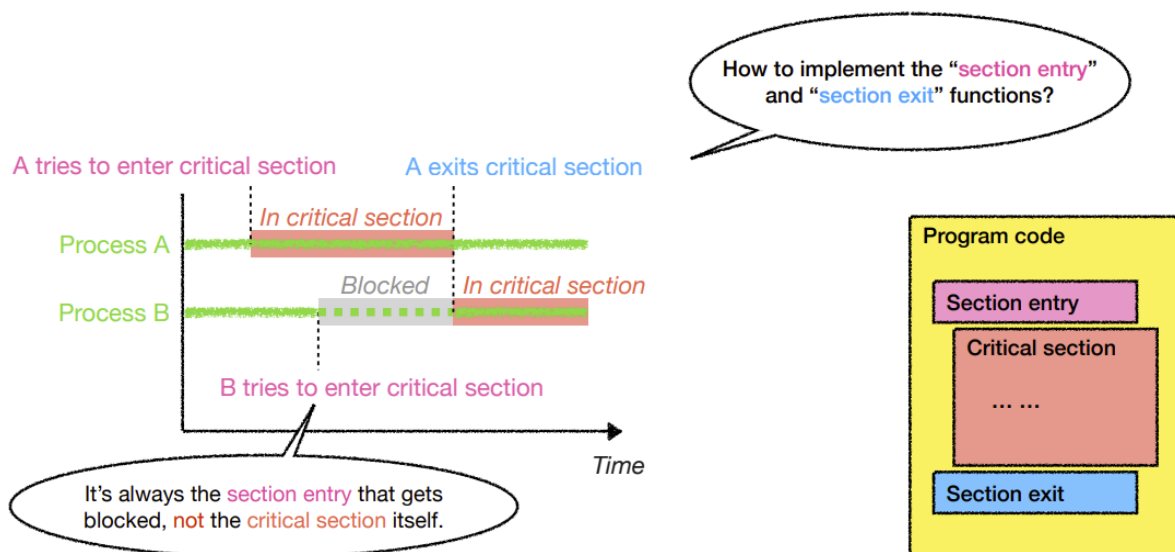
- The solution **cannot depend on the time spent inside the critical section** or assume the number of CPUs in the system.

**No process running outside its critical section may block other processes.**

- This ensures all processes can make **progress**. Otherwise, it may end up with a scenario where **all processes are blocked** but no process is inside its critical section.

**No process should have to wait forever to enter its critical section.**

- This guarantees **bounded waiting**, i.e., no process will **starve to death**.



- **Attempt #1: disabling interrupts**

When interrupts are disabled, **no context switch can occur**.

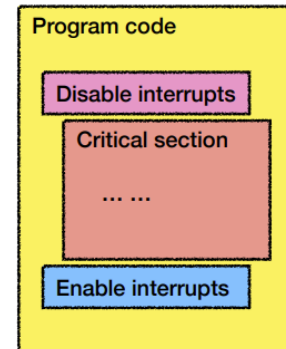
- When one process is in its critical section, the CPU will not be switched to another process.

**Is it correct?**

- Single-processor system: yes.
- **Multiprocessor system: NO!** Other CPUs can still access the shared resource.

**Is it a good solution?**

- Even on a single-processor system, it's a **terrible idea** to allow user processes to enable/disable interrupts at will.
- However, **within the OS kernel itself**, it's often convenient to disable interrupts for a few instructions.



- **Attempt #2: using a “lock” variable**

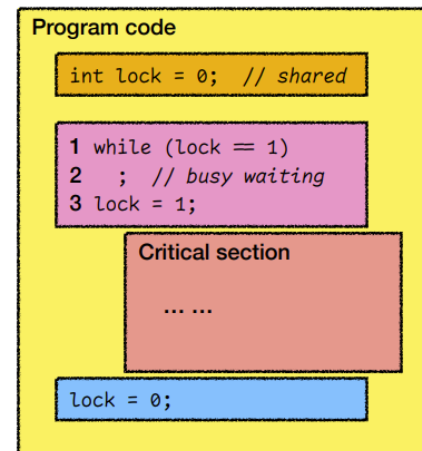
What about using a shared “lock” variable?

**Is it correct?**

- **NO!**
- It suffers from exactly the same **race condition**.

**Is it a good solution (even if it worked)?**

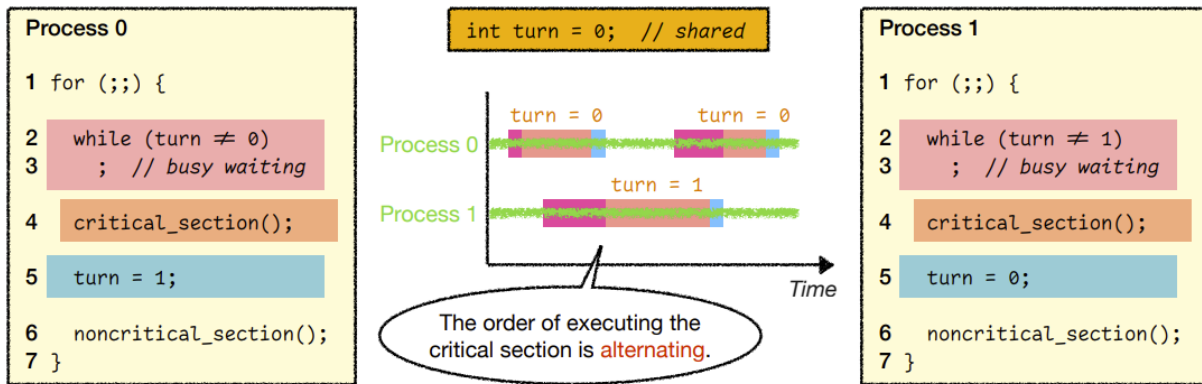
- **Busy waiting** (a.k.a. **spinning**) wastes CPU time.
- **Priority inversion problem**: a high-priority process may have to wait for a low-priority process to relinquish resource.



- **Busy Waiting**
  - The while loop keeps running and running again. It's also called spinning
- **Priority Inversion Problem**
  - a high-priority process may have to wait for a low-priority process to relinquish resource.

- **Attempt #3: strict alternation**

Use a shared variable to let processes **take turns** to enter the critical section.



Processes **take turns** to enter the critical section.

**Is it correct?**

- It does **avoid all races**.
- However...

**Is it a good solution?**

- It violates Requirement #3 of being a good solution:

**No process running outside its critical section may block other processes.**

#### • Attempt #4: Peterson's algorithm

## Critical sections

### Attempt #4: Peterson's algorithm

This is an improved version of strict alternation.

```

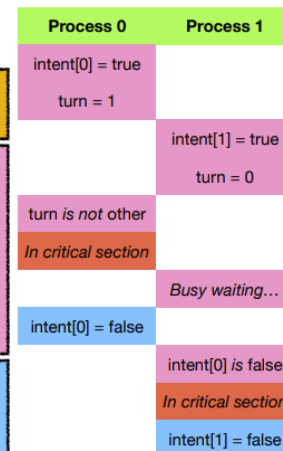
int intent[2] = {0, 0};           // who wants to enter critical section?
int turn;                         // whose turn is it?

1 void section_entry(int process) { // process is 0 or 1
2   int other = 1 - process;       // the other process is 1 or 0
3   intent[process] = true;        // I intend to enter critical section
4   turn = other;                  // make it the other process's turn
5   while (intent[other] && turn == other) // I'm a gentleman!
6     ;                            // busy waiting while it's not my turn
7 }

1 void section_exit(int process) {
2   intent[process] = false;       // simply undo my intent
3 }

```

Processes would **act as a gentleman**.  
"If you want to enter, I'll let you first."



- When the while loop runs, we can wait for the other process.
- When the while loop breaks, we can do our own work



Peterson's algorithm is a combination of using **lock variables** and **taking turns**.

### Is it correct?

- Yes on early hardware, where `intent[]` and turn propagate **immediately** and **atomically**.
- **Not anymore on modern hardware** due to relaxed memory consistency models.

### Is it a good solution?

- With a little **hardware support**, the solution could be much easier...

## • Attempt #5: spinlocks

Modern CPUs all have instructions that guarantee **atomic operation**.

The simplest one is a **test-and-set** instruction.

Conceptually, it behaves as if this code snippet is executed without interruption:

```
int test_and_set(int *ptr, int new) {  
    // this code executes atomically  
    int old = *ptr;  
    *ptr = new;  
    return old;  
}
```

- It returns the old value pointed to by `ptr` and simultaneously updates said value to `new`.

“

Atomic operation: The operation does a lot of things and can never be interrupted

```
int test_and_set(int *ptr, int new) {  
    // this code executes atomically  
    int old = *ptr;  
    *ptr = new;  
    return old;  
}
```

How to implement `section_entry()` and `section_exit()` using the atomic `test_and_set` instruction?

```
int lock = 0; // 0: available, 1: held
```

```
void section_entry(int *lock) {  
    while (test_and_set(lock, 1) == 1)  
        ; // busy waiting  
}
```

```
void section_exit(int *lock) {  
    *lock = 0;  
}
```

- We can use the `test_and_set()` to set the `lock` to 1 and check its old value.
  - If the value is already 1, it means someone else is in critical section, and we have to wait.
  - If it's 0, then we can enter the critical section

**Spinlocks** are built upon CPU instructions that guarantee **atomic operation**.

- **Example:** test-and-set, compare-and-swap, fetch-and-add, load-linked/store-conditional.

### Is it correct?

- Yes on a **preemptive** scheduler. (Beware of the **priority inversion problem**!)

### Is it a good solution?

- **Simplicity:** **yes!**
- **Fairness:** no guarantees. A process may spin forever under **contention**, leading to **starvation**.
- **Performance:**
  - Single-processor system: bad. With  $N$  processes contending for a lock,  $N-1/N$  of CPU time is wasted.
  - Multiprocessor system: effective if [# of processes]  $\approx$  [# of CPUs] and **the critical section is short**.

“

If a scheduler is nonpreemptive, nothing can work

“

Priority Inversion problem exists in this solution

### - Attempt #5.1: too much spinning?

```
int test_and_set(int *ptr, int new) {  
    // this code executes atomically  
    int old = *ptr;  
    *ptr = new;  
    return old;  
}
```

Can you change one thing to make it less wasteful?

```
int lock = 0; // 0: available, 1: held
```

```
void section_entry(int *lock) {  
    while (test_and_set(lock, 1) == 1)  
        sched_yield(); // give up the CPU  
}
```

```
void section_exit(int *lock) {  
    *lock = 0;  
}
```

- When the lock is 1, instead of waiting forever, we can just give up the CPU to others

While better than spinning, this **run-and-yield** approach is still costly:

- Before the process holding the lock gets to run again, every other process still needs to run and yield.
- **Context switching** is expensive.

What if you can **block** instead of spinning?

It's time to introduce a new OS primitive.



- **Attempt #6: semaphores**

A **semaphore** is an object with a **non-negative** integer value.


- The value **must be initialized** before being used.

There are two operations: `down()` and `up()`.


- See “[man 7 sem\\_overview](#)” for more details.

```
void down(semaphore *sem) {
    // atomic operation
    if (*sem > 0) {
        *sem = *sem - 1;
    } else {
        block on sem;
    }
}
```

```
void up(semaphore *sem) {
    // atomic operation
    if (some process is blocked on sem) {
        let one such process proceed
    } else {
        *sem = *sem + 1;
    }
}
```



```
void down(semaphore *sem) {
    // atomic operation
    if (*sem > 0) {
        *sem = *sem - 1;
    } else {
        block on sem;
    }
}
```



```
void up(semaphore *sem) {
    // atomic operation
    if (some process is blocked on sem) {
        let one such process proceed
    } else {
        *sem = *sem + 1;
    }
}
```

## Program code

```
// init a binary semaphore
semaphore sem = 1;
```

```
down(&sem);
```

### Critical section

□ □ □ □ □ □

```
up(&sem);
```

“

Binary semaphore: the value of it is either 0 or 1

- `up()`
  - If there're more than one process blocked, it can choose any of the process to proceed.
  - there's no specific order

We just used a **binary semaphore** to implement a **mutex** (“**mutual exclusion**”).

### Is it correct?

- Yes!

### Is it a good solution?

- Yes!

Actually, semaphores are more powerful than guaranteeing **mutual exclusion**.

- Semaphores can be used to realize **process synchronization**, i.e., **to coordinate** the set of processes so they can produce meaningful output.
- If the value of a semaphore is  $> 1$ , we call it a **counting semaphore**.
- Starvation?
  - Possible

### • Summary

- Today we introduced **6 different attempts for critical sections**
- The first three attempts didn't actually work
- **Peterson's algorithm** was a good attempt back in the days, but now it doesn't work anymore
- **Spinlocks** and **Semaphores** are the two approaches that are widely used today
  - **Spinlocks** are mainly used for short critical sections, as it uses busy waiting
  - **Semaphores** put the processes to the block state instead of busy waiting, so that it doesn't waste CPU power