

# OS: Lecture 11

---

## OS: Lecture 11

Review: Critical Sections

    Requirements

    Attempts

        Attempt 5

        Attempt 6

    Classical IPC problems

    The producer-consumer problem

        Requirements

        Implementation using semaphores

            If we swapped lines 6 & 7 in the producer...

            deadlock

        Summary

    The dining philosopher problem

        Modeling

        Requirement #1: mutual exclusion

            Attempt #1: mutex for each chopstick

        Requirement #2: synchronization

            Attempt #2: mutex for each chopstick with backoff

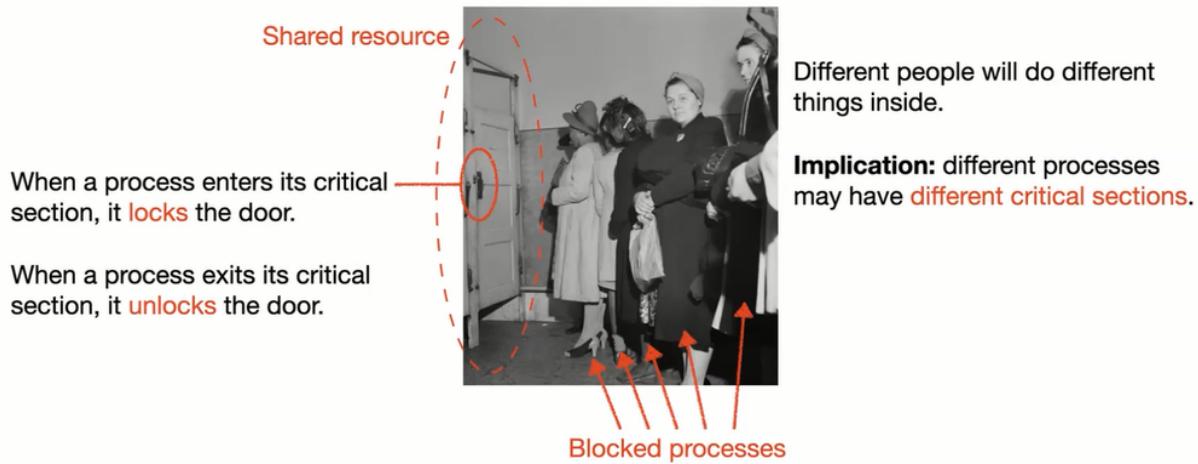
            Attempt #3: mutex for each chopstick with random backoff

            Attempt #4: global mutex

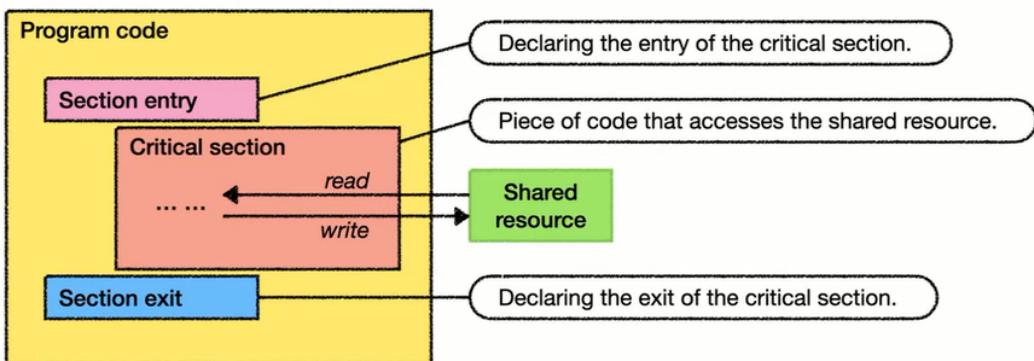
            Attempt #5: one semaphore per philosopher

        Summary

## Review: Critical Sections



A **critical section** (a.k.a. **critical region**) is a **piece of code** that accesses a shared resource.



- Requirements

## Requirements

### No two processes may be simultaneously inside their critical sections.

- This is the **mutual exclusion** requirement: when one process is inside its critical section, any attempts to go inside the critical sections by other processes are **not allowed**.

### No assumptions may be made about the speeds or the number of CPUs.

- The solution **cannot depend on the time spent inside the critical section** or assume the number of CPUs in the system.

### No process running outside its critical section may block other processes.

- This ensures all processes can make **progress**. Otherwise, it may end up with a scenario where **all processes are blocked** but no process is inside its critical section.

### No process should have to wait forever to enter its critical section.

- This guarantees **bounded waiting**, i.e., no process will **starve to death**.

- Attempts

Attempt #1: disabling interrupts.

Attempt #2: using a “lock” variable.

Attempt #3: strict alternation.

Attempt #4: Peterson’s algorithm.

Attempt #5: spinlocks.

Attempt #6: semaphores.

- Attempt 5

```
int test_and_set(int *ptr, int new) {
    // this code executes atomically
    int old = *ptr;
    *ptr = new;
    return old;
}
```

How to implement `section_entry()` and `section_exit()` using the atomic `test_and_set` instruction?

```
int lock = 0; // 0: available, 1: held
```

```
void section_entry(int *lock) {
    while (test_and_set(lock, 1) == 1)
        ; // busy waiting
}
```

```
void section_exit(int *lock) {
    *lock = 0;
}
```

- It's a good solution with some limitations
  - The critical section has to be small
  - The number of processes should not be more than the number of cores in your system

- Attempt 6

A **semaphore** is an object with a **non-negative** integer value.

- The value **must be initialized** before being used.

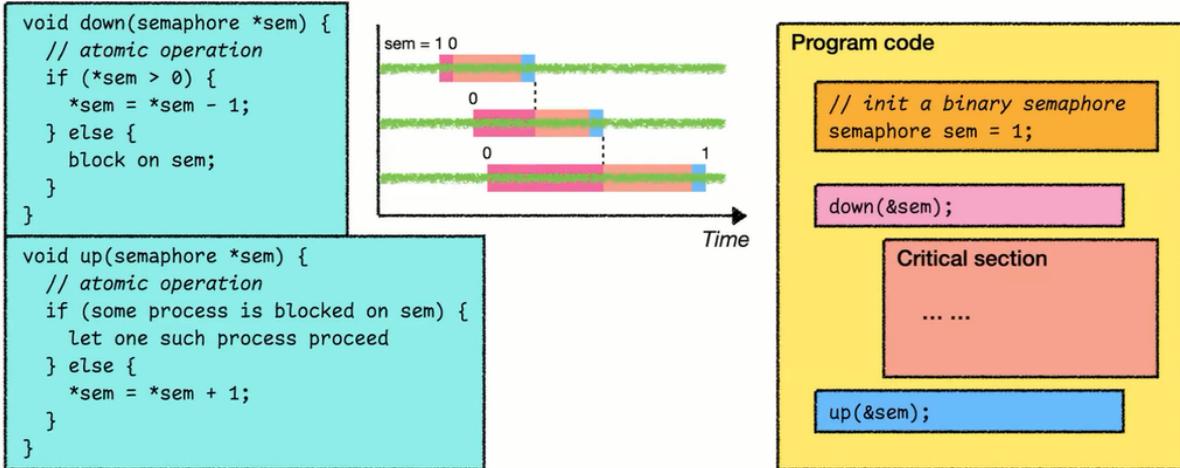
There are two operations: `down()` and `up()`.

- See “`man 7 sem_overview`” for more details.

```
void down(semaphore *sem) {
    // atomic operation
    if (*sem > 0) {
        *sem = *sem - 1;
    } else {
        block on sem;
    }
}
```

```
void up(semaphore *sem) {
    // atomic operation
    if (some process is blocked on sem) {
        let one such process proceed
    } else {
        *sem = *sem + 1;
    }
}
```





## Classical IPC problems

### The producer-consumer problem

- It models access to a bounded buffer.

### The dining philosopher problem

- It models processes competing for exclusive access to a limited number of resources (e.g., I/O devices).

### The readers and writers problem

- It models access to a database.

### The sleeping barber problem

- It models a queueing situation (*please stay on the line...*).

## The producer-consumer problem

### The producer-consumer problem

- It models access to a bounded buffer.

### The dining philosopher problem

- It models processes competing for exclusive access to a limited number of resources (e.g., I/O devices).

### The readers and writers problem

- It models access to a database.

### The sleeping barber problem

- It models a queueing situation (*please stay on the line...*).

## a.k.a. the bounded-buffer problem

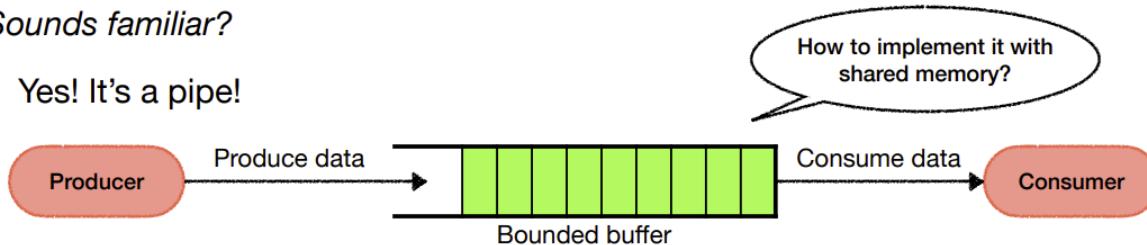
Two processes share a fixed-size buffer.

The **producer** inserts data to the **tail** of the buffer.

The **consumer** removes data from the **head** of the buffer.

*Sounds familiar?*

- Yes! It's a pipe!



- Requirements

### Synchronization requirement #1

When the **producer** wants to insert an item into the buffer, but **the buffer is already full...**

What should the producer and the consumer do?

- The **producer** should **block**.
- The **consumer** should **wake up** the producer **after it has consumed an item**.

### Synchronization requirement #2

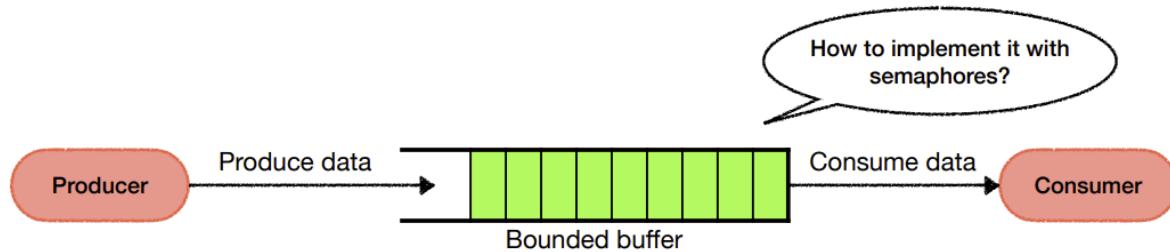
When the **consumer** wants to remove an item from the buffer, but **the buffer is empty...**

What should the producer and the consumer do?

- The **consumer** should **block**.
- The **producer** should **wake up** the consumer **after it has produced an item**.

Of course, we also need **mutual exclusion**...

- No two processes can access the **shared** buffer at the same time.



- Implementation using semaphores

```
semaphore mutex = 1;           // controls access to critical section
semaphore empty = BUFFER_SIZE; // counts empty buffer slots
semaphore filled = 0;          // counts filled buffer slots
```

Why do we need three semaphores?

```
1 void producer() {
2     int item;
3
4     for (;;) {
5         item = produce_item();
6         down(&empty);    // decrement empty count
7         down(&mutex);   // enter critical section
8         insert(item);   // put new item in buffer
9         up(&mutex);    // exit critical section
10        up(&filled);  // increment filled count
11    }
12 }
```

```
1 void consumer() {
2     int item;
3
4     for (;;) {
5         down(&filled);    // decrement filled count
6         down(&mutex);   // enter critical section
7         item = remove(); // take item from buffer
8         up(&mutex);    // exit critical section
9         up(&empty);    // increment empty count
10        consume_item(item);
11    }
12 }
```

- Only `insert()` and `remove()` are between the `down(mutex)` and `up(mutex)`,
  - as we want our critical sections to be as short as possible
- `empty` counts the number of empty buffer slots
  - `down(&empty)`, if the buffer is already full and `empty = 0`, the **producer** is blocked
  - `up(&empty)` means that the **consumer** has removed something from the buffer and there's one more empty slot
    - If the **producer** was blocked, the consumer will wake up the producer by adding the `empty`
    - If it's not blocked, it means we still have plenty of empty space here
- `filled` counts the number of filled buffer slot, initial = 0
-

`mutex` guarantees mutual exclusion.

`empty` and `filled` are for synchronization.

- `empty` represents the number of empty slots.
  - The **producer blocks** when this value is 0, i.e., there is no empty slot, i.e., the buffer is already full.
- `filled` represents the number of filled slots.
  - The **consumer blocks** when this value is 0, i.e., there is no filled slot, i.e., the buffer is empty.

**Invariant:** `empty + filled = buffer size.`

```
semaphore mutex = 1;           // controls access to critical section
semaphore empty = BUFFER_SIZE; // counts empty buffer slots
semaphore filled = 0;          // counts filled buffer slots
```

Can we swap lines 6 & 7 in the producer?

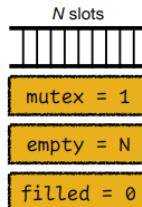
```
1 void producer() {
2   int item;
3
4   for (;;) {
5     item = produce_item();
6     down(&empty); // decrement empty count
7     down(&mutex); // enter critical section
8     insert(item); // put new item in buffer
9     up(&mutex); // exit critical section
10    up(&filled); // increment filled count
11  }
12 }
```

```
1 void consumer() {
2   int item;
3
4   for (;;) {
5     down(&filled); // decrement filled count
6     down(&mutex); // enter critical section
7     item = remove(); // take item from buffer
8     up(&mutex); // exit critical section
9     up(&empty); // increment empty count
10    consume_item(item);
11  }
12 }
```

- If we swapped lines 6 & 7 in the producer...

## If we swapped lines 6 & 7 in the producer...

Producer  
Consumer



```
1 void producer() {
2   int item;
3
4   for (;;) {
5     item = produce_item();
6     down(&mutex); // enter critical section
7     down(&empty); // decrement empty count
8     insert(item); // put new item in buffer
9     up(&mutex); // exit critical section
10    up(&filled); // increment filled count
11  }
12 }
```

```
1 void consumer() {
2   int item;
3
4   for (;;) {
5     down(&filled); // decrement filled count
6     down(&mutex); // enter critical section
7     item = remove(); // take item from buffer
8     up(&mutex); // exit critical section
9     up(&empty); // increment empty count
10    consume_item(item);
11  }
12 }
```

Producer

Lines 4-11 for  $N$  times

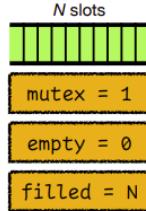
Consumer

```

1 void producer() {
2     int item;
3
4     for (;;) {
5         item = produce_item();
6         down(&mutex); // enter critical section
7         down(&empty); // decrement empty count
8         insert(item); // put new item in buffer
9         up(&mutex); // exit critical section
10        up(&filled); // increment filled count
11    }
12 }
```

```

1 void consumer() {
2     int item;
3
4     for (;;) {
5         down(&filled); // decrement filled count
6         down(&mutex); // enter critical section
7         item = remove(); // take item from buffer
8         up(&mutex); // exit critical section
9         up(&empty); // increment empty count
10        consume_item(item);
11    }
12 }
```



Producer

Lines 4-11 for  $N$  times

Line 5

Line 7

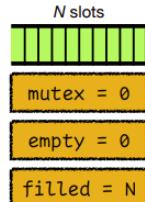
Consumer

```

1 void producer() {
2     int item;
3
4     for (;;) {
5         item = produce_item();
6         down(&mutex); // enter critical section
7         down(&empty); // decrement empty count
8         insert(item); // put new item in buffer
9         up(&mutex); // exit critical section
10        up(&filled); // increment filled count
11    }
12 }
```

```

1 void consumer() {
2     int item;
3
4     for (;;) {
5         down(&filled); // decrement filled count
6         down(&mutex); // enter critical section
7         item = remove(); // take item from buffer
8         up(&mutex); // exit critical section
9         up(&empty); // increment empty count
10        consume_item(item);
11    }
12 }
```



Producer

Lines 4-11 for  $N$  times

Line 5

Line 7

Line 6 Blocked on empty

Consumer

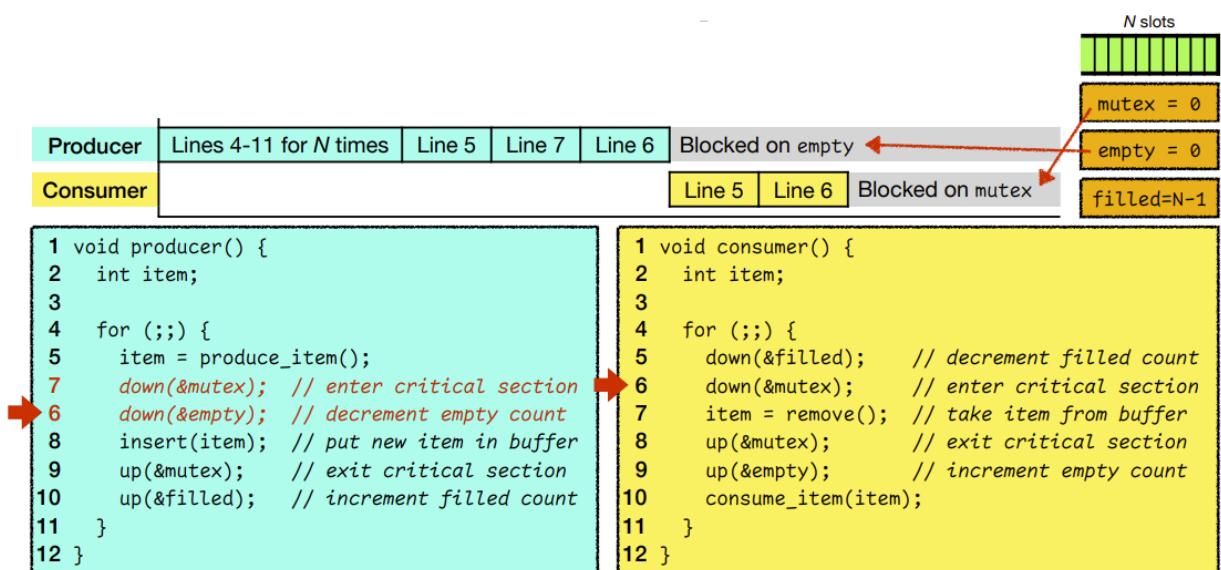
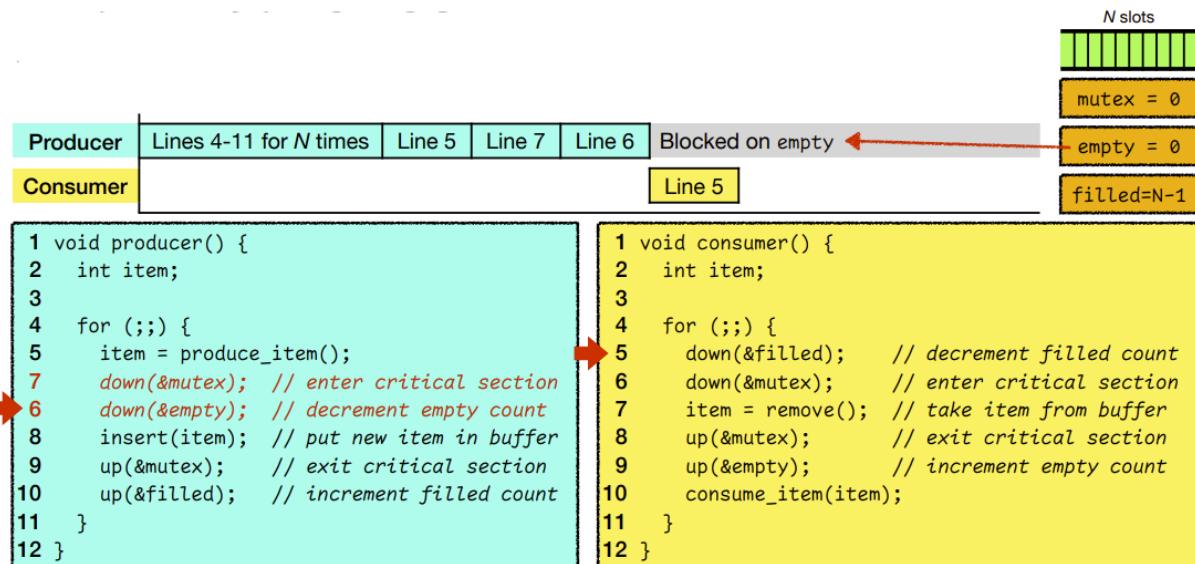
```

1 void producer() {
2     int item;
3
4     for (;;) {
5         item = produce_item();
6         down(&mutex); // enter critical section
7         down(&empty); // decrement empty count
8         insert(item); // put new item in buffer
9         up(&mutex); // exit critical section
10        up(&filled); // increment filled count
11    }
12 }
```

```

1 void consumer() {
2     int item;
3
4     for (;;) {
5         down(&filled); // decrement filled count
6         down(&mutex); // enter critical section
7         item = remove(); // take item from buffer
8         up(&mutex); // exit critical section
9         up(&empty); // increment empty count
10        consume_item(item);
11    }
12 }
```

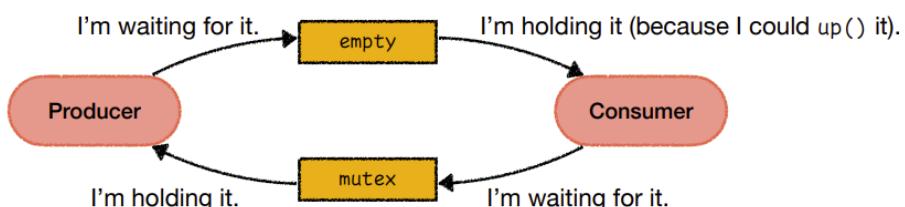




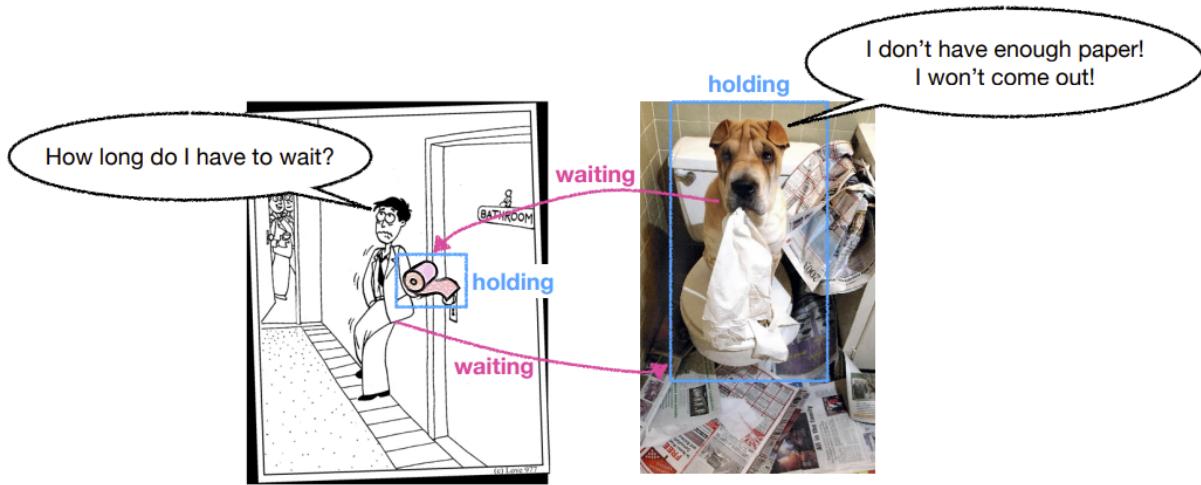
This scenario is called a **deadlock**.

It happens when there is a **circular wait**...

- The **producer** is waiting for the consumer to `up()` the `empty` semaphore.
- The **consumer** is waiting for the producer to `up()` the `mutex` semaphore.



- **deadlock**



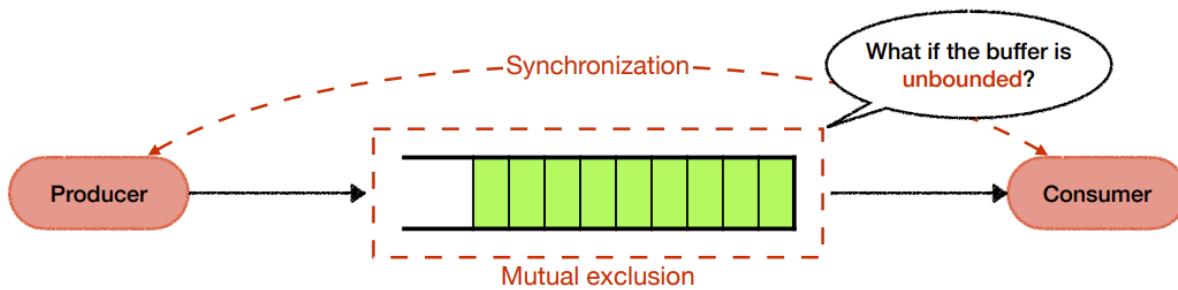
- **Summary**

### Mutual exclusion

A **binary semaphore** is used as a lock to protect the critical sections.

### Synchronization

Two are used to monitor the status of the buffer. counting semaphoresoblem



## The dining philosopher problem

### The producer-consumer problem

- It models access to a bounded buffer.

### The dining philosopher problem

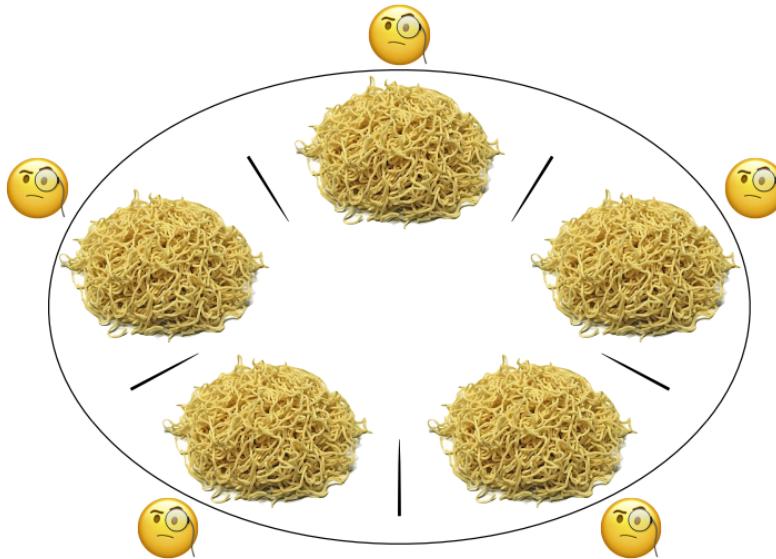
- It models processes competing for exclusive access to a limited number of resources (e.g., I/O devices).

### The readers and writers problem

- It models access to a database.

### The sleeping barber problem

- It models a queueing situation (please stay on the line...).



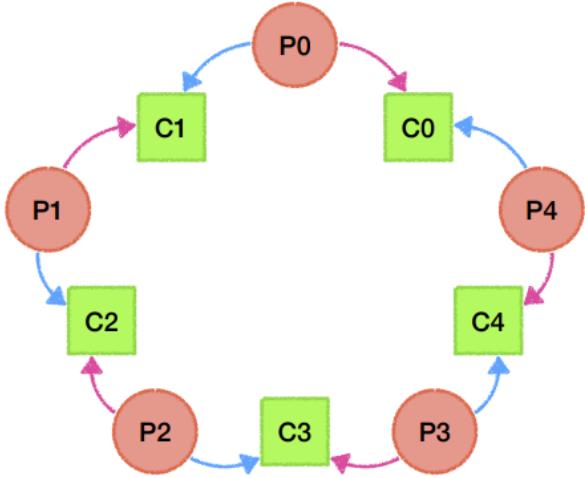
5 philosophers around a table;  
5 instant ramen noodles;  
5 chopsticks in between.

A philosopher does only two things in their entire life:  
**eating** and **thinking**.

In order to eat, a philosopher needs **exactly two chopsticks**.

How to design a protocol so that they can enjoy their life?

- Modeling



For each philosopher  $i$ ,

- *left chopstick* =  $i$
- *right chopstick* =  $(i + 1) \% N$
- *LEFT neighbor* =  $(i + N - 1) \% N$
- *RIGHT neighbor* =  $(i + 1) \% N$

States:

- THINKING
- HUNGRY ← Trying to get chopsticks.
- EATING ← Critical section.

- Requirement #1: mutual exclusion

### Requirement #1: mutual exclusion

*What if there is no mutual exclusion?*

- Then, while you're eating, the two philosophers sitting next to you will **steal your chopsticks!**

- Attempt #1: mutex for each chopstick

## Attempt #1: mutex for each chopstick

```

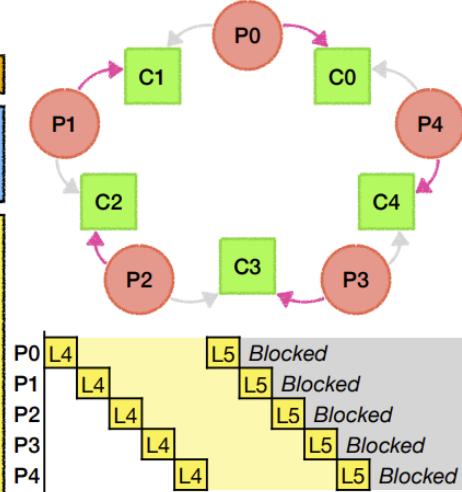
semaphore chopsticks[N]; // what are the initial values?

void take_chopstick(int i) {
    down(&chopsticks[i]);
}

void put_chopstick(int i) {
    up(&chopsticks[i]);
}

1 void philosopher(int i) {
2     for (;;) {
3         think();
4         take_chopstick(i);
5         take_chopstick((i + 1) % N);
6         eat();                                // critical section
7         put_chopstick((i + 1) % N);
8         put_chopstick(i);
9     }
10 }

```



- Requirement #2: synchronization

## Requirement #2: synchronization

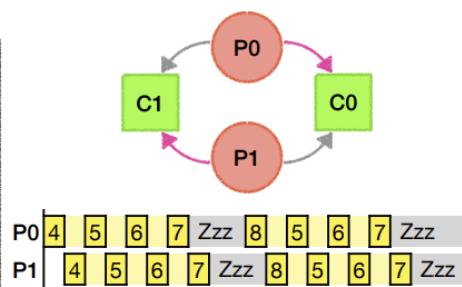
The solution should avoid any potential **deadlock** or **starvation** scenarios.

- Attempt #2: mutex for each chopstick with backoff

```

1 void philosopher(int i) {
2     for (;;) {
3         think();
4         take_chopstick(i);
5         while (try_to_take_chopstick((i + 1) % N) == FAILED) {
6             /* it either succeeds and takes the chopstick,
7                or fails and does not take the chopstick. */
8             put_chopstick(i);
9             sleep(1);           // wait for 1s and try again
10            take_chopstick(i);
11        }
12        eat();                  // critical section
13        put_chopstick((i + 1) % N);
14        put_chopstick(i);
15    }
16 }

```



This scenario is called a **livelock**.

No process is blocked, but they don't make any progress.

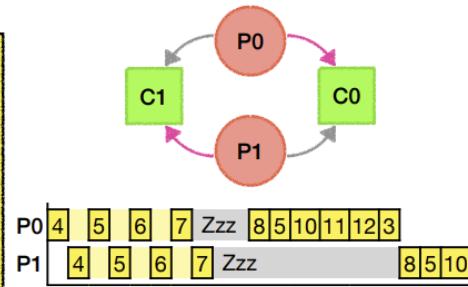
It's a kind of **starvation**, literally.

- Attempt #3: mutex for each chopstick with random backoff

```

1 void philosopher(int i) {
2   for (;;) {
3     think();
4     take_chopstick(i);
5     while (try_to_take_chopstick((i + 1) % N) == FAILED) {
6       /* it either succeeds and takes the chopstick,
7          or fails and does not take the chopstick. */
8       put_chopstick(i);
9       sleep(random()); // wait a random time
10      take_chopstick(i);
11    }
12    eat();           // critical section
13    put_chopstick((i + 1) % N);
14    put_chopstick(i);
15  }
16 }

```



This is good in most applications.

In fact, many network protocols are designed this way.

Is there a non-randomized solution?

### - Attempt #4: global mutex

```

semaphore mutex = 1;
semaphore chopsticks[N]; // what are the initial values?

```

```

1 void philosopher(int i) {
2   for (;;) {
3     think();
4     down(&mutex);
5     take_chopstick(i);
6     take_chopstick((i + 1) % N);
7     eat();           // critical section
8     put_chopstick((i + 1) % N);
9     put_chopstick(i);
10    up(&mutex);
11  }
12 }

```

### Is it correct?

- Any deadlock? — No.
- Any starvation? — No.

This solution is theoretically correct.

### Is it a good solution?

With the global mutex, there is no parallelization.

Only one philosopher can be eating at any time.

The performance is very bad.

### - Attempt #5: one semaphore per philosopher

Before we propose the final solution, let's see what issues we've been facing...

- Modeling each chopstick as a semaphore is intuitive, but it's not working.
- We are afraid to `down()` a chopstick, as that may lead to a deadlock.

However, our real issue is that **philosophers** need to eat, **not chopsticks**.

So, instead of worrying about chopsticks, we should guarantee...

- When a **philosopher** is eating, their **left and right neighbors** cannot eat.

## Attempt #5: one semaphore per philosopher

```

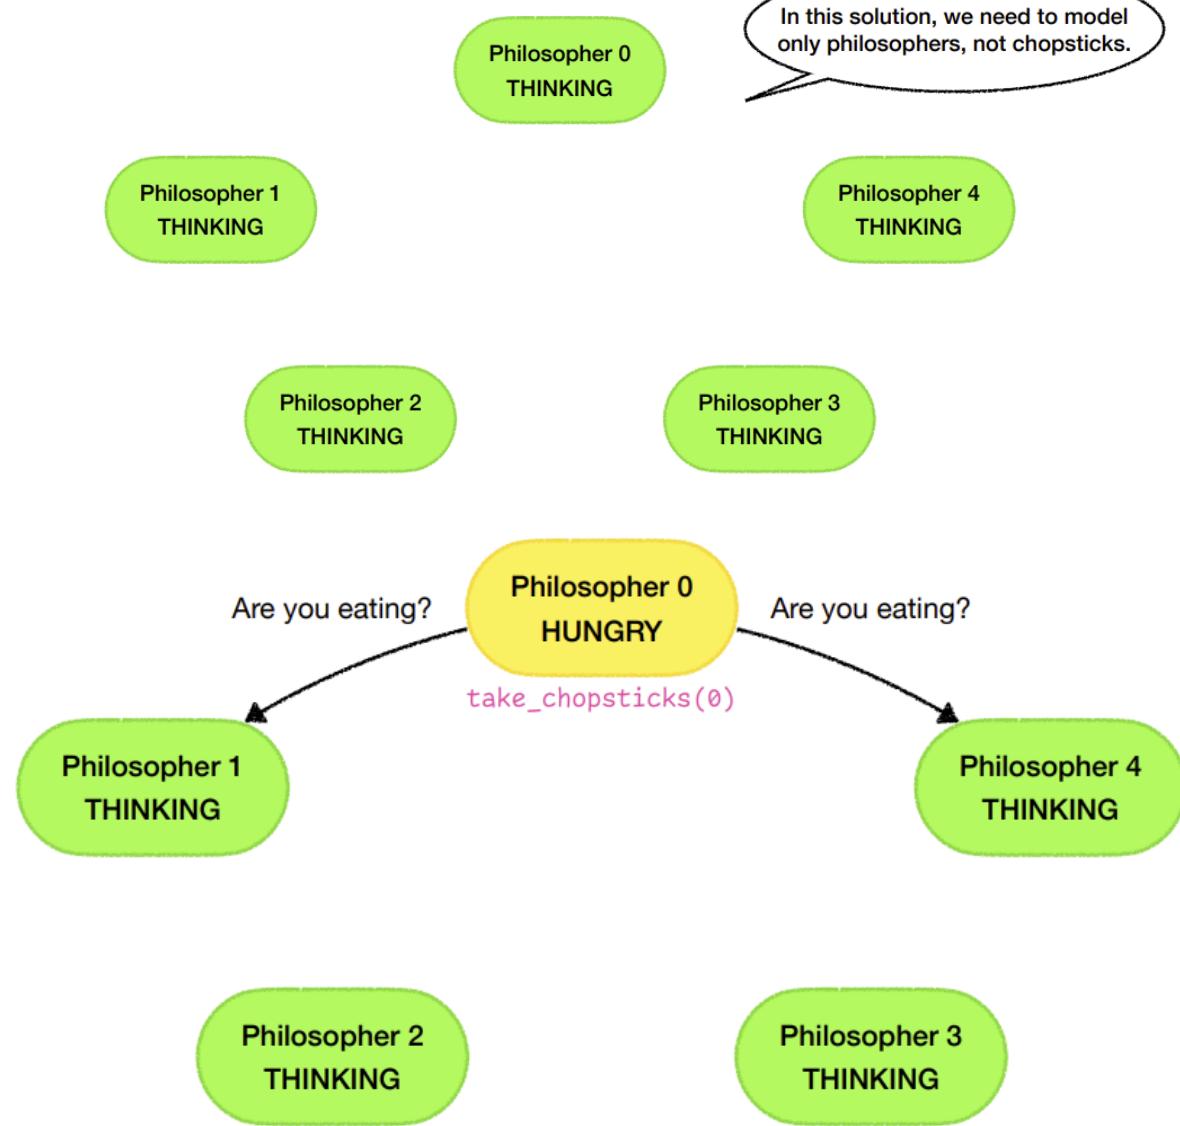
int state[N];           // THINKING, HUNGRY, or EATING
semaphore mutex = 1;
semaphore sem[N];       // what are the initial values?

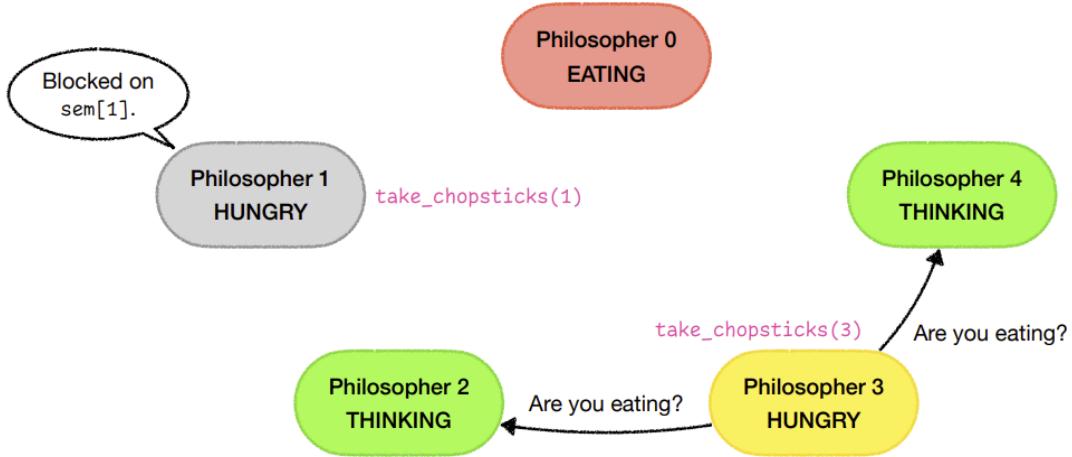
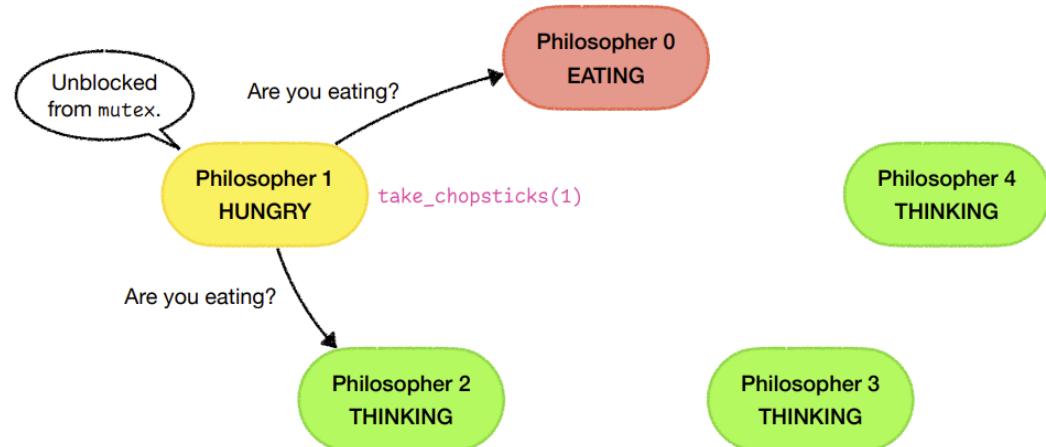
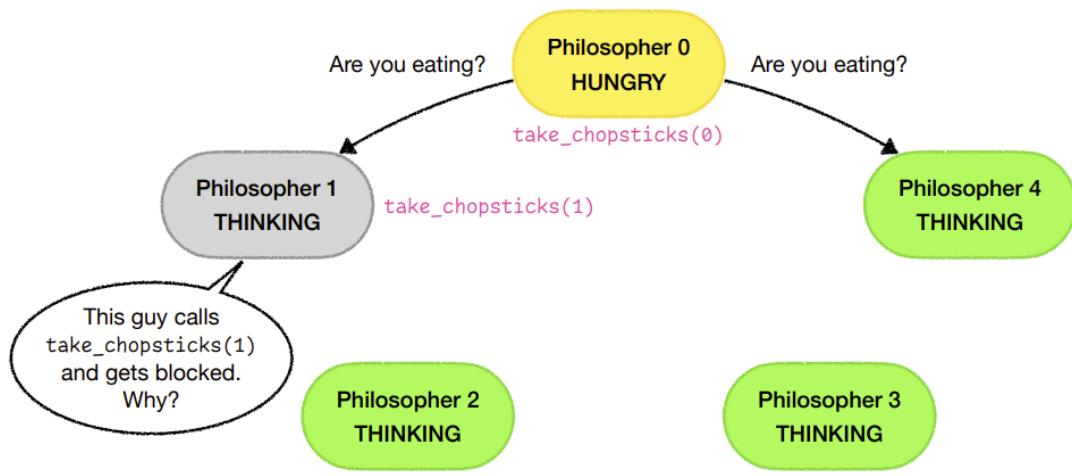
void take_chopsticks(int i) {
    down(&mutex);
    state[i] = HUNGRY;
    test(i);          // try to take both chopsticks
    up(&mutex);
    down(&sem[i]);   // block if cannot take chopsticks
}

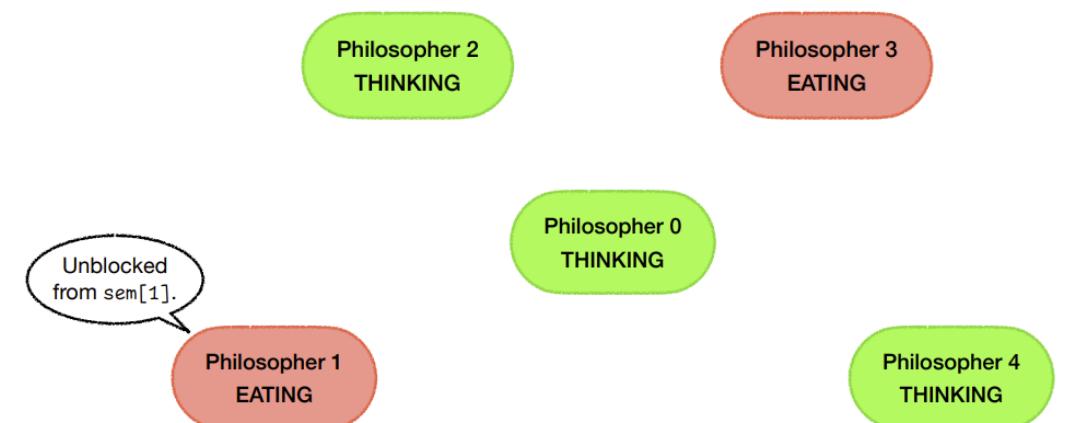
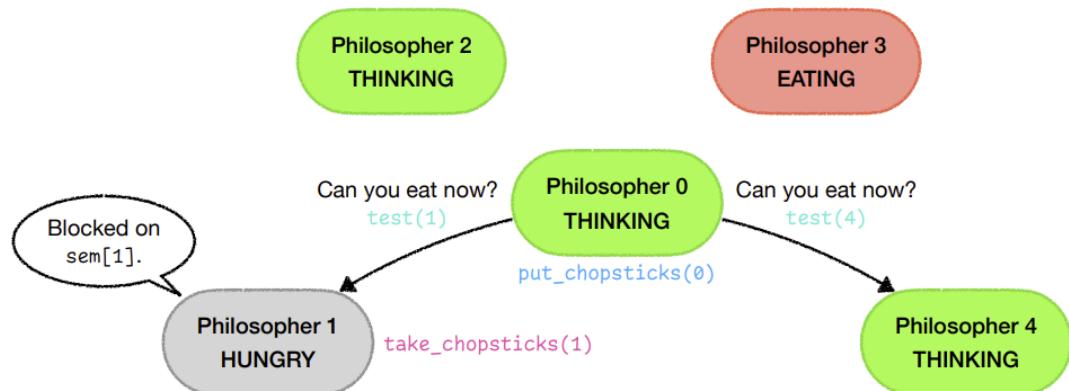
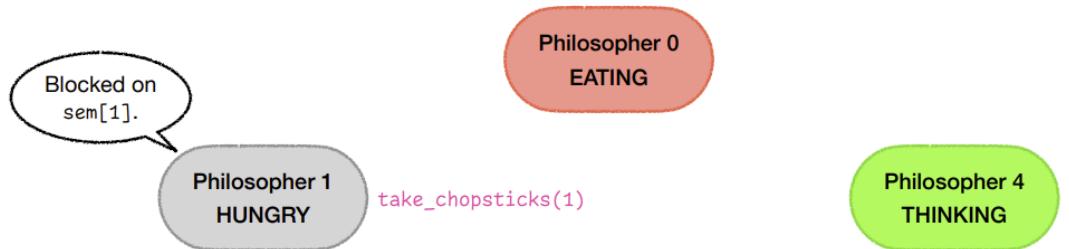
void put_chopsticks(int i) {
    down(&mutex);
    state[i] = THINKING;
    test(LEFT);        // see if left neighbor can now eat
    test(RIGHT);       // see if right neighbor can now eat
    up(&mutex);
}

void philosopher(int i) {
    for (;;) {
        think();
        take_chopsticks(i); // section entry
        // block until I take both chopsticks
        eat();              // critical section
        put_chopsticks(i); // section exit
    }
}

```







- Summary

This problem is once again about **mutual exclusion** and **synchronization**.

Multiple **processes** compete for multiple **resources**.

Each process needs **two distinct resources** to enter the critical section.

The **random backoff** approach is adopted in many network protocols.

The **final solution** is **deadlock-free**, **starvation-free**, and allows the maximum **parallelism** for any number of processes.