

OS: Lecture 12

OS: Lecture 12

Review

The producer-consumer problem

The dining philosopher problem

The readers and writers problem

Accessing a database

Requirements

Attempt #1: mutex for the database

Attempt #2: allow concurrent readers

The sleeping barber problem

The client-server model

The solution

IPC problems Summary

Threads: lightweight processes

What's a thread?

Why threads?

Thread models

Where to implement threads?

Many-to-one model

One-to-one model

Many-to-many model


Midterm

Review


- The producer-consumer problem
- The dining philosopher problem

The readers and writers problem

- Accessing a database



Multiple processes are allowed to **read** the database at the same time.



Multiple processes are allowed to **read** the database at the same time.

If a process is **writing** the database, no other processes may have access to the database.

How to program the **readers** and **writers**?

- Requirements

Mutual exclusion

- The database is a **shared resource**.

Synchronization

- When a **reader** is reading, other **readers** are allowed to read the database.
- When a **reader** is reading, no **writers** are allowed to write the database.
- When a **writer** is writing, no **readers** or **writers** are allowed to access the database.

Concurrency

- **Concurrent** access from multiple **readers** should be allowed.

- Attempt #1: mutex for the database

```
semaphore db = 1;
```

```
void reader() {  
    for (;;) {  
        down(&db);           // section entry  
        data = read_database(); // critical section  
        up(&db);             // section exit  
        consume_data(data);  
    }  
}
```

```
void writer() {  
    for (;;) {  
        data = produce_data();  
        down(&db);           // section entry  
        write_database(data); // critical section  
        up(&db);            // section exit  
    }  
}
```

No two processes can be in their critical sections at the same time.

However, how to allow **concurrent** access from multiple **readers**?

- **Attempt #2: allow concurrent readers**

Multiple **readers** do not need to mutually exclude one another.

However, as long as there is a **reader** in the system, we need to exclude **writers**.

Any ideas?

The **first reader** that comes to the system should **down(&db)**.

The **last reader** that exits from the system should **up(&db)**.

We just need to **count the number of readers** in the system.

```
semaphore db = 1;           // controls access to the database  
semaphore mutex = 1;       // controls access to "reader_count"  
int reader_count = 0;      // # of processes reading or wanting to read
```

```
void reader() {  
    for (;;) {  
        down(&mutex);  
        if (++reader_count == 1)  
            down(&db); // the first reader locks db  
        up(&mutex);  
        data = read_database(); // critical section  
        down(&mutex);  
        if (--reader_count == 0)  
            up(&db); // the last reader unlocks db  
        up(&mutex);  
        consume_data(data);  
    }  
}
```

```
void writer() {  
    for (;;) {  
        data = produce_data();  
        down(&db);           // locks db  
        write_database(data); // critical section  
        up(&db);            // unlocks db  
    }  
}
```


This solution meets all our requirements.

However, **it gives readers a higher priority than writers**.

As long as there is a steady supply of **readers**, **writers** will suffer from **starvation**.

The sleeping barber problem

The barber shop



1 barber, 1 barber chair;
 n chairs for waiting customers.

If there are no customers, the barber falls asleep.

When a customer arrives, they wake up the barber.

If more customers arrive while the barber is busy...

- They wait as long as there are empty chairs; or
- They leave the shop if all chairs are full.

How to program the barber and the customers?

- The client-server model

This problem is similar to various **queueing** situations in a **client-server model**.

What are the processes?

- One **long-running server** process: the barber;
- Many **transient client** processes: the customers.

Any shared resources?

- The **server** (barber) itself is shared by all clients (customers). Only **one client** can be served at a time.
- **Limited wait queue**: n chairs. When the queue is full, new clients will be dropped.

We need **mutual exclusion** and proper **synchronization**.

- The solution

```
semaphore barbers = 0; // # of barbers ready to cut hair (why 0 instead of 1?)
semaphore customers = 0; // # of customers waiting for service (not being cut)
int waiting = 0; // same as above (because there's no way to read a semaphore's value)
semaphore mutex = 1; // controls access to "waiting"
```

```
void barber() {
    for (;;) {
        down(&customers); // sleep if no customer
        down(&mutex);
        --waiting;
        up(&barbers); // a barber is ready to serve
        up(&mutex);
        cut_hair(); // outside critical section
    }
}
```

```
void customer() {
    down(&mutex);
    if (waiting < NUM_CHAIRS) {
        ++waiting;
        up(&customers); // wake up barber if needed
        up(&mutex);
        down(&barbers); // wait if no barber is free
        get_haircut(); // outside critical section
    } else {
        up(&mutex); // leave if the shop is full
    }
}
```

It's **data-race-free** and **deadlock-free**. Any **starvation**?

IPC problems involve...

- One or more **shared resources**;
- **Multiple processes** that must be synchronized.

A good solution need to...

- Guarantee **mutual exclusion**;
- Guarantee proper **synchronization** among processes;
- Be **deadlock-free**;
- Be **starvation-free**.

Threads: lightweight processes

- What's a thread?

A thread is an **execution entity within a process**.

So far, we have only discussed **single-threaded** processes.

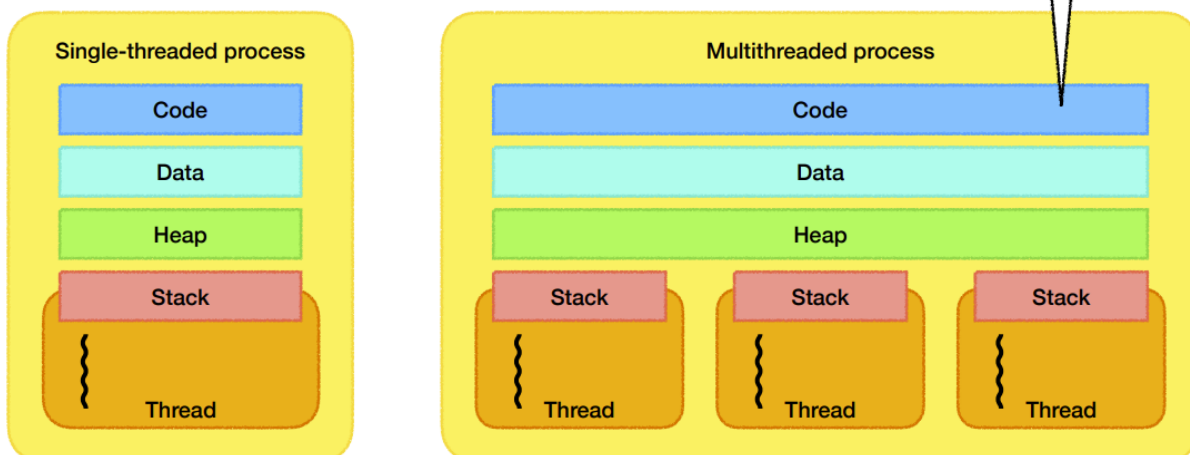
A **multithreaded process** can have more than one execution in it.

Example: a word processor may have...

- A thread interacting with the user;
- A thread formatting text in the background;
- A thread handling automatic backups.

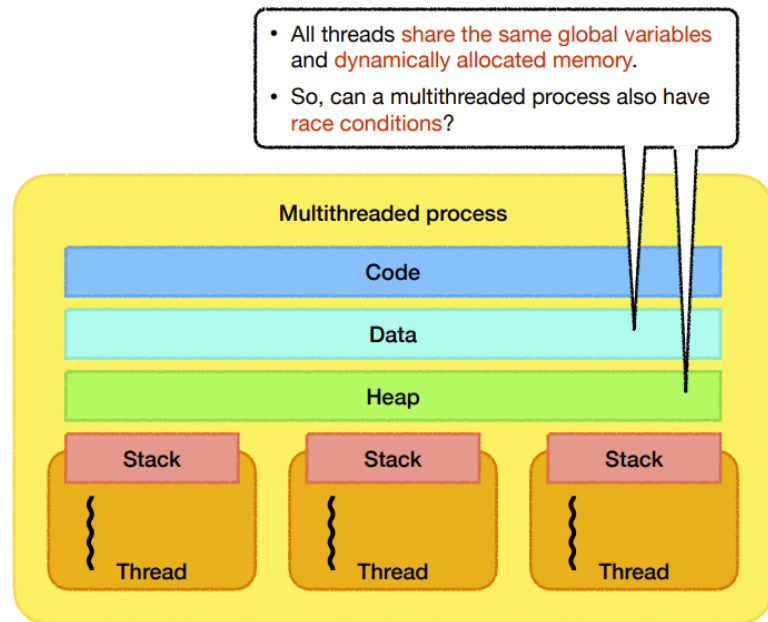

Threads

What's a thread?



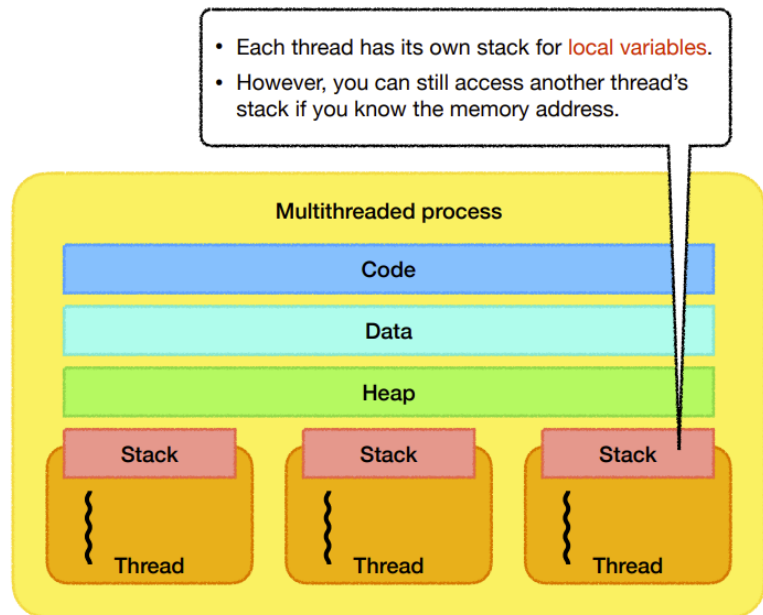

Threads

What's a thread?

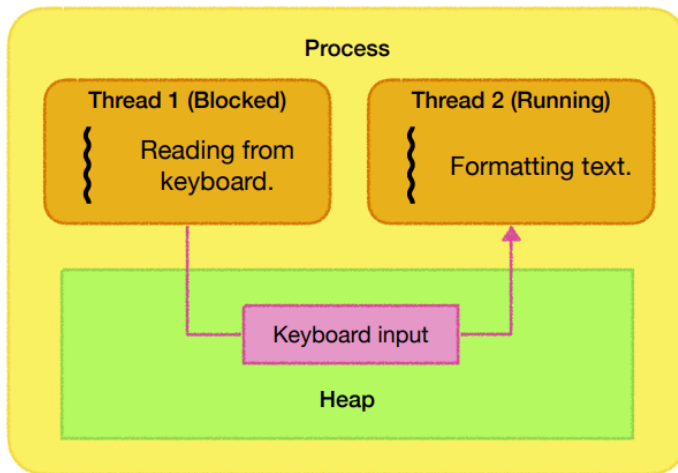


Threads

What's a thread?



- Why threads?



It allows **multitasking** within a process:

- One thread waits for the user input;
 - Another thread performs computation.
- ⇒ Better **performance** and **responsiveness**.

Threads are **easier to create and destroy** than processes (10-100× faster).

Threads share the **same address space**; so sharing data is easy.

Thread models

• Where to implement threads?

Many-to-one model

- Implement threads in user space.

One-to-one model

- Implement threads in the kernel.

Many-to-many model

- Hybrid implementations.

• Many-to-one model

The threads library is entirely in **user space**.


- The kernel knows nothing about threads (“**green threads**”).
- Implemented in earlier operating systems.

Pros:

- Does not need OS support.
- Fast (no trap, no context switch, no need to flush cache...).

Cons:

- When a **blocking system call** is invoked, all threads will be blocked.
- **Page faults** (discussed later) in a thread will block the entire process.
- **No preemption** of threads due to the absence of clock interrupts.



- One-to-one model

The kernel manages all the threads.


- Each thread is mapped to a **kernel thread**.
- Implemented in most modern operating systems.

Pros:

- When a thread blocks, the kernel can **schedule another thread** (from either the same process or a different process).

Cons:

- Creating and destroying threads are **more expensive**.
- What happens when a multithreaded process **forks**?
- When a **signal** comes in, which thread should handle it?



- Many-to-many model

The threads library **multiplexes** user-level threads onto kernel threads.

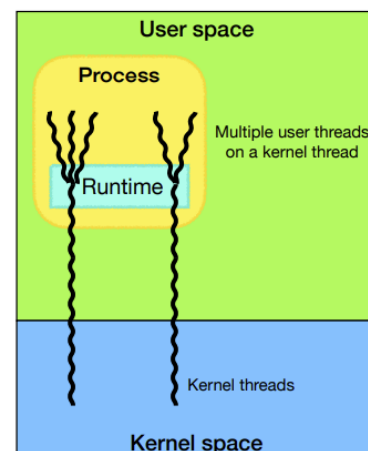
- It's a hybrid, **two-level** model.
- Implemented in many language runtimes (Erlang, Go, Haskell, JVM...).

Pros:

- The best of both worlds (more flexible).
- No restrictions on the number of threads.

Cons:

- More complex to implement.



Midterm

When: Thursday, March 9, from 12:30 p.m. to 1:45 p.m.

Where: Brightspace Quizzes.

What: Everything up to today's class.

- All lectures.
- Lab 1 (nyuc) and Lab 2 (nyush).
- Homework 1 and 2.

You must **connect to the Zoom session** and **turn on your webcam** for the entire duration of the exam.

You can refer to...

- Slides.
- Textbooks.
- Docker & CIMS servers (e.g., read man pages, use gcc to try out programs).
- Online websites (you can search existing information, but **you are not allowed to ask questions and seek answers**).

However, keep in mind that by doing these things, you will likely **waste a lot of time** with little gain.

You must not communicate with any other people or AI-based assistants.

Format

Multiple-choice questions?

All questions are “**choose one out of two**” questions!

- True or False **or I don't know!**
- Possible or Impossible **or I don't know!**

No fussing, no guessing, just **choose “I don't know” to get 50% points!**

Don't leave any questions blank!